

CSCI3150 Introduction to Operating Systems

Lecture 4: Threads

Hong Xu

<https://github.com/henryhxu/CSCI3150>

Processes

- ❑ Recall that a process includes many things
 - ◆ An address space (all the code and data)
 - ◆ OS resources (e.g., open files) and accounting information
 - ◆ Execution state (PC, SP, regs, etc.)
- ❑ Creating a new process is costly because of all the data structures that must be allocated and initialized
 - ◆ Recall struct proc in Solaris
 - ◆ ...which does not even include page tables, perhaps TLB flushing, etc.
- ❑ Communication between processes is costly because it goes through the OS
 - ◆ Overhead of system calls and copying data

Parallel Programs

- ❑ To execute these programs we need to
 - ◆ Create several processes that execute in parallel
 - ◆ Cause each to map to the same address space to share data
 - They are all part of the same computation
 - ◆ Have the OS schedule these processes in parallel
- ❑ This situation is **very inefficient**
 - ◆ **Space**: PCB, page tables, etc.
 - ◆ **Time**: create data structures, fork and copy addr space, etc.
- ❑ Is it possible to have more **efficient**, yet **cooperative** “processes”?

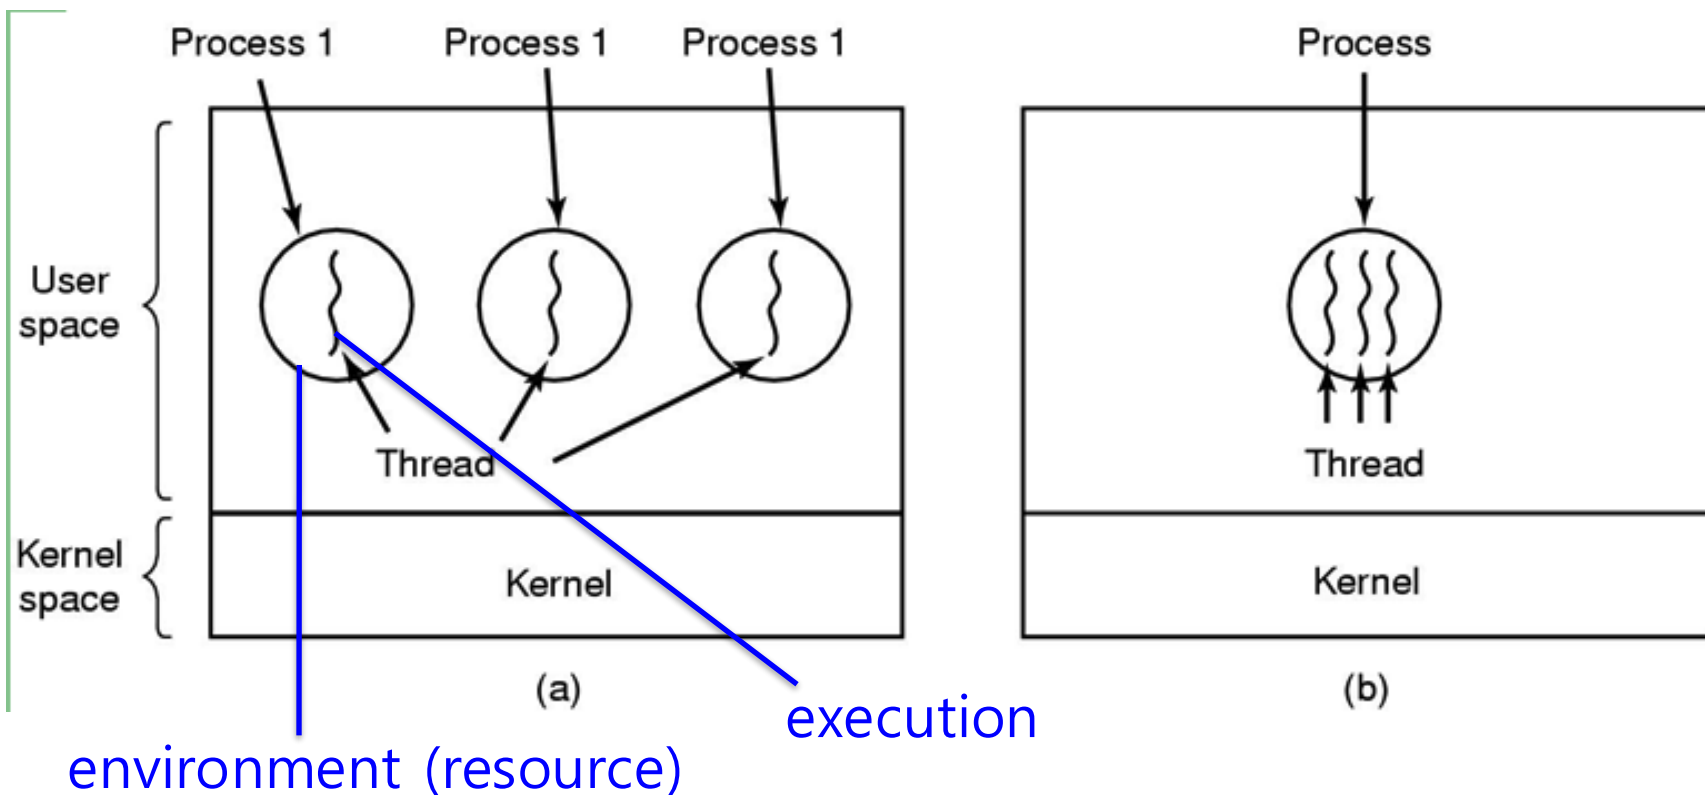
Rethinking Processes

- ❑ What is similar in these cooperating processes?
 - ◆ They all share the same code and data (address space)
 - ◆ They all share the same privileges
 - ◆ They all share the same resources (files, sockets, etc.)
- ❑ What don't they share?
 - ◆ Each has its own execution state: PC, SP, and registers
- ❑ **Key idea:** Why don't we **decouple** the concept of a process from its execution state?
 - ◆ **Process:** address space, privileges, resources, etc.
 - ◆ **Execution state:** PC, SP, registers
- ❑ Exec state also called **thread of control**, or **thread**

Threads

- ▣ Modern OSes (Windows, modern Unix) separate the concepts of processes and threads
 - ◆ The **thread** defines a sequential execution stream within a process (PC, SP, registers)
 - ◆ The **process** defines the address space and general process attributes
- ▣ A thread is bound to a single process
 - ◆ Processes, however, can have **multiple** threads
- ▣ Threads become the unit of scheduling
 - ◆ Processes are now the **containers** in which threads execute

Threads: lightweight processes



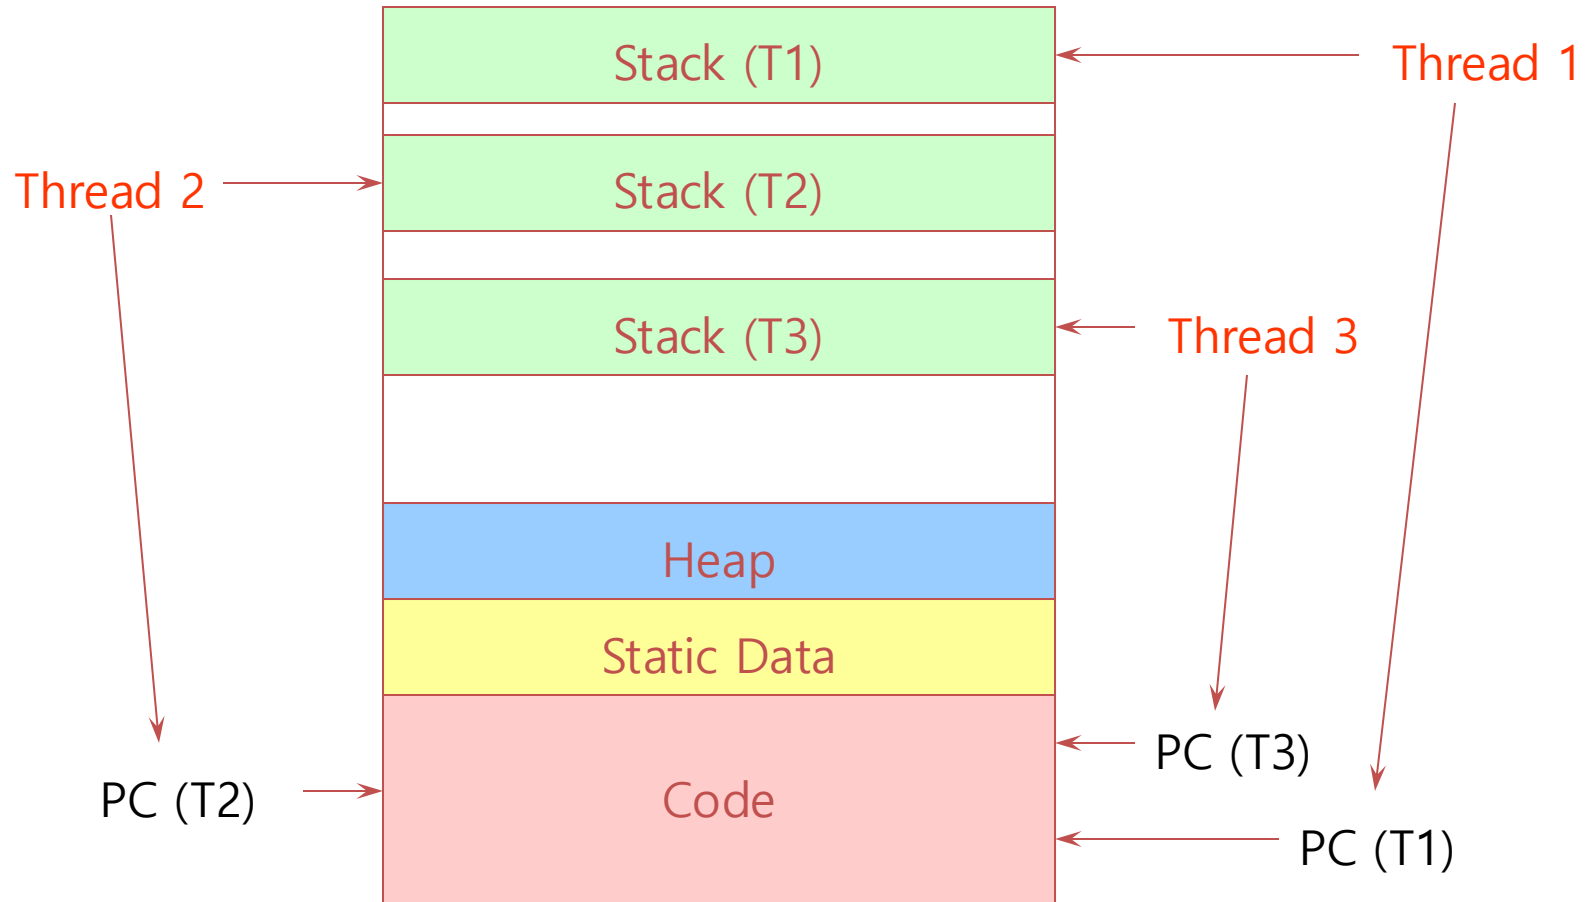
(a) Three processes each with one thread

(b) One process with three threads

The thread model

- ▣ Shared information
 - ◆ Processor info: parent process, time, etc
 - ◆ Memory: segments, page table, and stats, etc
 - ◆ I/O and file: communication ports, directories and file descriptors, etc.
- ▣ Private state
 - ◆ State (ready, running and blocked)
 - ◆ Registers
 - ◆ Program counter
 - ◆ Execution stack
- ▣ Each thread executes independently

Threads in a Process



Threads: Concurrent Servers

- ▣ Using `fork()` to create new processes to handle requests in parallel is an overkill for such a simple task
- ▣ Recall our forking Web server:

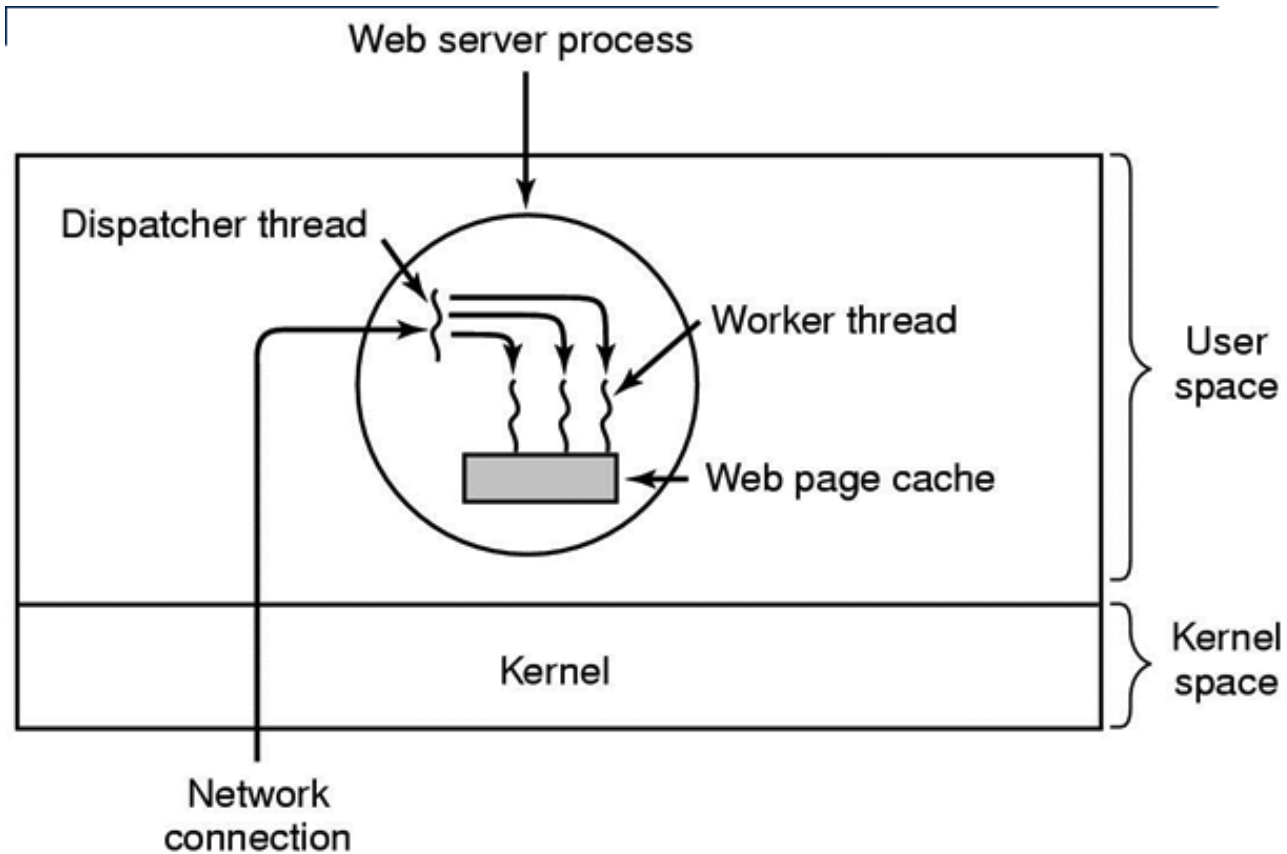
```
while (1) {  
    int sock = accept();  
    if ((child_pid = fork()) == 0) {  
        //Handle client request  
        //Close socket and exit  
    } else {  
        //Close socket  
    }  
}
```

Threads: Concurrent Servers


- ▣ Instead, we can create a new thread for each request

```
web_server() {  
    while (1) {  
        int sock = accept();  
        thread_create(handle_request, sock);  
    }  
}  
  
handle_request(int sock) {  
    //Process request  
    close(sock);  
}
```

Thread usage: web server



Thread usage: word processor



- A thread can wait for I/O, while the other threads can still run.
- What if it is single-threaded?

Kernel-Level Threads

- ▣ We have taken the execution aspect of a process out and separated it into threads
 - ◆ To make concurrency cheaper
- ▣ As such, the OS now manages threads *and* processes
 - ◆ All thread operations are implemented in the kernel
 - ◆ The OS schedules all threads in the system
- ▣ OS-managed threads are called **kernel-level threads** or **lightweight processes**
 - ◆ Windows: **threads**
 - ◆ Solaris: **lightweight processes (LWP)**
 - ◆ POSIX Threads (pthreads): **PTHREAD_SCOPE_SYSTEM**

Kernel-Level Thread Limitations

- ❑ Kernel-level threads make concurrency much cheaper than processes
 - ◆ Much less state to allocate and initialize
- ❑ However, for fine-grained concurrency, kernel-level threads still suffer from too much overhead
 - ◆ Thread operations still require system calls
 - Ideally, want thread operations to be **as fast as a procedure call**
- ❑ For such fine-grained concurrency, need even “cheaper” threads

User-Level Threads

- ❑ To make threads fast, they need to be implemented at user level
 - ◆ Kernel-level threads are managed by the OS
 - ◆ User-level threads are managed by the run-time system (user-level library)
- ❑ User-level threads are small, cheap, and fast
 - ◆ A thread is simply represented by a PC, registers, stack, and small thread control block (TCB)
 - ◆ Creating a new thread, switching between threads, and synchronizing threads are done via procedure calls
 - No kernel involvement
 - ◆ User-level thread operations 100x faster than kernel threads
 - ◆ pthreads: PTHREAD_SCOPE_PROCESS

User-Level Thread Limitations

- ▣ Yet, user-level threads are not a perfect solution
 - ◆ As with everything else, they are a tradeoff
- ▣ User-level threads are **invisible** to the OS
 - ◆ They are not well integrated with the OS
- ▣ As a result, the OS can make poor decisions
 - ◆ Scheduling a process with idle threads
 - ◆ **Blocking a process whose thread initiated an I/O, even though the process has other threads that can execute**
- ▣ Solving this requires communication between the kernel and the user-level thread manager

Kernel- vs. User-level Threads

- ▣ Kernel-level threads
 - ◆ Integrated with OS (informed scheduling)
 - ◆ Slow to create, manipulate, synchronize
- ▣ User-level threads
 - ◆ Fast to create, manipulate, synchronize
 - ◆ Not integrated with OS (uninformed scheduling)
- ▣ Understanding the differences between kernel- and user-level threads is important
 - ◆ For programming (correctness, performance)
 - ◆ For test-taking

Kernel- vs. User-level Threads

- ▣ Or use **both** kernel- and user-level threads
 - ◆ Can associate a user-level thread with a kernel-level thread
 - ◆ Or, multiplex user-level threads on top of kernel-level threads
- ▣ Golang today uses user-level threads
 - ◆ Multiplex multiple Goroutines (user-level threads) on multiple kernel level threads

Implementing Threads

- ▣ Implementing threads has a number of issues
 - ◆ Interface
 - ◆ Context switch
 - ◆ Preemptive vs. Non-preemptive
 - What do they mean?
 - ◆ Scheduling
 - ◆ Synchronization (next lecture)
- ▣ Focus on user-level threads
 - ◆ Kernel-level threads are similar to original process management and implementation in the OS

Sample Thread Interface

- `thread_create`(procedure_t, arg)
 - ◆ Create a new thread of control
 - ◆ Start executing procedure_t
- `thread_yield`()
 - ◆ Voluntarily give up the processor
- `thread_exit`()
 - ◆ Terminate the calling thread; also thread_destroy
- `thread_join`(target_thread)
 - ◆ Suspend the execution of calling thread until target_thread terminates

Thread Scheduling

- ▣ For user-level thread: scheduling occurs entirely in user-space
- ▣ The thread scheduler determines when a thread runs
- ▣ It uses queues to keep track of what threads are doing
 - ◆ Just like the OS and processes
 - ◆ But it is implemented at user-level in a library
- ▣ Run queue: Threads currently running (usually one)
- ▣ Ready queue: Threads ready to run
- ▣ Are there wait queues?

Review of threads

- ▣ What are shared among threads of the same process? What are not?
 - ◆ Why cannot they share the same stack?
 - ◆ How threads of the same process communicate with each other?
- ▣ Trade-off between kernel level threads and user level threads
- ▣ Blocking system call
 - ◆ Blocking system call: an I/O system call that will wait for the I/O to complete before returning
- ▣ How do we implement user-level threads

Non-Preemptive Scheduling

- Threads voluntarily give up the CPU with `thread_yield`

Ping Thread

```
while (1) {  
    printf("ping\n");  
    thread_yield();  
}
```

Pong Thread

```
while (1) {  
    printf("pong\n");  
    thread_yield();  
}
```

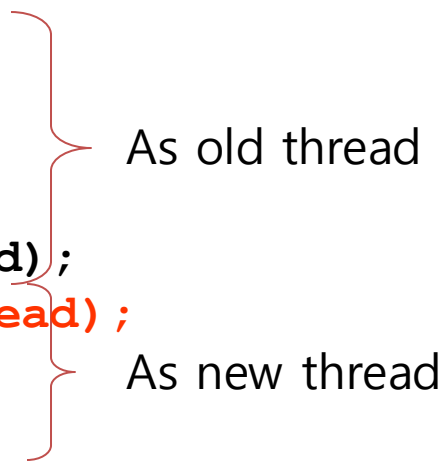
- What is the output of running these two threads?

thread_yield()

- ❑ Wait a second. How does thread_yield() work?
 - ◆ it gives up the CPU to another thread
 - ◆ In other words, it **context switches** to another thread
- ❑ So what does it mean for thread_yield to return?
 - ◆ It means that *another thread* called thread_yield!
- ❑ Execution trace of ping/pong
 - ◆ `printf("ping\n");`
 - ◆ `thread_yield();`
 - ◆ `printf("pong\n");`
 - ◆ `thread_yield();`
 - ◆ ...

Implementing thread_yield()

```
thread_yield() {  
    thread_t old_thread = current_thread;  
    current_thread = get_next_thread();  
    append_to_queue(ready_queue, old_thread);  
    context_switch(old_thread, current_thread);  
    return;  
}
```



As old thread

As new thread

- ▣ The magic step is invoking context_switch()
- ▣ Why do we need to call append_to_queue()?

Thread Context Switch

- ▣ The context switch routine does all the magic
 - ◆ Saves context of the currently running thread (`old_thread`)
 - Push all machine state onto its stack (*except stack pointer*)
 - ◆ Restores context of the next thread
 - Pop all machine state from the next thread's stack
 - ◆ The next thread becomes the current thread
 - ◆ Return to caller as new thread
- ▣ This is all done in assembly language

Wait a minute

- ❑ Non-preemptive threads have to voluntarily give up CPU
 - ◆ Only voluntary calls to `thread_yield()`, or `thread_exit()` causes a context switch
- ❑ What if a thread never release the CPU (never calls `thread_yield()`)?
- ❑ We need preemptive user-level thread scheduling

Preemptive Scheduling

- ▣ Preemptive scheduling causes an involuntary context switch
 - ◆ Need to regain control of processor asynchronously
- ▣ How?
 - ◆ *Use timer interrupt*
 - ◆ Timer interrupt handler forces current thread to “call” `thread_yield`
 - How?

Process vs. thread

- ❑ Multithreading is only an option for “cooperative tasks”
 - ◆ Trust and sharing
- ❑ Process
 - ◆ Strong isolation but poor performance
- ❑ Thread
 - ◆ Good performance but share too much
- ❑ Example: web browsers
 - ◆ Safari: multithreading
 - one webpage can crash entire Safari
 - ◆ Google Chrome: each tab has its own process

Summary

- ▣ The operating system as a large multithreaded program
 - ◆ Each process executes as a thread within the OS
- ▣ Multithreading is also very useful for applications
 - ◆ Efficient multithreading requires fast primitives
 - ◆ Processes are too heavyweight
- ▣ Solution is to separate threads from processes
 - ◆ Kernel-level threads much better, but still significant overhead
 - ◆ User-level threads even better, but not well integrated with OS

Summary cont.

- But the problem with threads is that, we need to explicitly manage concurrency
 - ◆ Why? Access to shared data
 - ◆ Synchronization, in the next lecture...
 - ◆ It's inherently hard, very hard
- For this and many other reasons, some people argue threads should not be used until absolutely necessary for performance
 - ◆ Check out the optional reading by Prof. John Outsterhaut in the 90s...