

CSCI3150 Introduction to Operating Systems


Lecture 5: Synchronization 1: Critical Region and Lock

Hong Xu

<https://github.com/henryhxu/CSCI3150>

Synchronization: why?

- ▣ A running computer has multiple processes; each process may have multiple threads



- ▣ Threads **access the shared data!**
 - ◆ Need proper sequencing
- ▣ Analogy: two people talking at the same time

A simple game

- ▣ Two volunteers to play two threads
 - ◆ Producer: produce 1 cookie per iteration
 - Step1: increment the counter on the board
 - Step2: put one cookie on the table
 - ◆ Consumer:
 - Step1: read out the counter LOUD
 - Step2a: if the counter is zero, go back to step1
 - Step2b: if the counter is nonzero, take a cookie from the table
 - Step 3: decrement counter on the board
 - ◆ Rule: only one should "operate" at any time
- ▣ You are the OS, decide when to context switch
 - ◆ Can you get them into "trouble" before cookies run out?

A simple game (cont.)

- ◆ Producer: produce 1 cookie per iteration
 - Step1: increment the counter on the board
 - Step2: put one cookie on the table
 - ◆ Consumer:
 - Step1: read out the counter LOUD
 - Step2a: if the counter is zero, go back to step1
 - Step2b: if the counter is nonzero, take a cookie from the table
 - Step 3: decrement counter on the board
- Switch to consumer, what will happen?*
- Switch to producer, what will happen?*

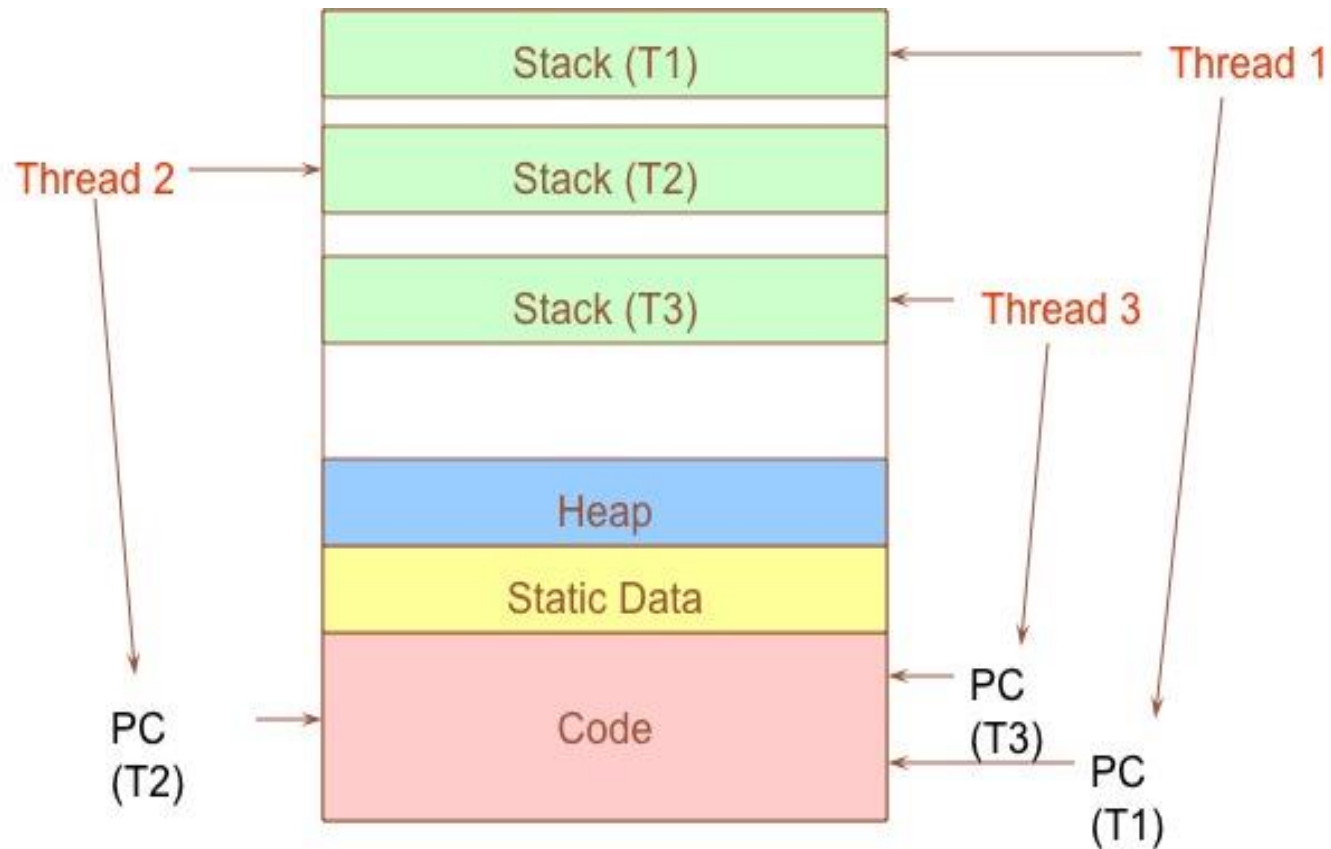
Data races

- Why are we having this problem?
- Reason:
 - ◆ concurrency
 - ◆ data sharing
- What are shared in this game?
 - ◆ Share the counter
 - ◆ Share the cookie
- A *race* occurs when *correctness* of the program depends on one thread reaching point x before another thread reaches point y

Shared Resources

- ▣ The problem is that two concurrent threads (or processes) access **shared resources** without **synchronization** (i.e. coordination)
 - ◆ Known as a **race condition** (memorize this)
- ▣ We need mechanisms to control access to shared resources in face of concurrency
 - ◆ So we can reason about how the program will operate
- ▣ **Shared data structures**
 - ◆ Buffers, queues, lists, hash tables, etc.

When are resources shared?



When are resources shared?

- ❑ Local variables are **not shared** (private)
 - ◆ Live on the stack
 - ◆ Each thread has its own stack
 - ◆ **Never pass/share/store a pointer to a local variable on the stack for thread T1 to another thread T2**
- ❑ Global variables and static objects are **shared**
 - ◆ Stored in the data segment, accessible by any thread
- ❑ Dynamic objects and other heap objects are **shared**
 - ◆ Allocated from heap with malloc/free or new/delete

A Classic Example

- ▣ Suppose we have to implement a function to handle withdrawals from a bank account:

```
withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return amount;  
}
```

- ▣ Now suppose that you and your significant other share a bank account with a balance of \$1000.
- ▣ Then you each go to separate ATM machines and simultaneously withdraw \$100 from the account.

Example Continued

- ▣ We'll represent the situation by creating a separate thread for each person to do the withdrawals
- ▣ These threads run on the same bank server:

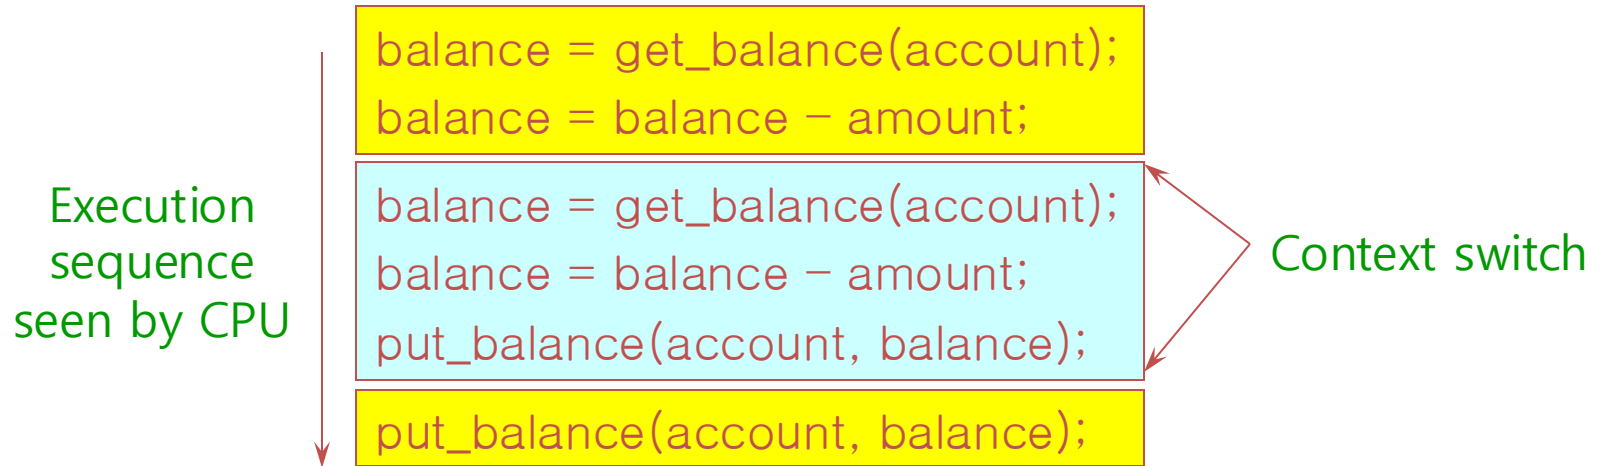
```
withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return amount;  
}
```

```
withdraw (account, amount) {  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return amount;  
}
```

- ▣ What's the problem with this implementation?
 - ◆ Think about possible schedules of these two threads

Interleaved Schedules

- Execution of the two threads can be interleaved:



- What is the balance of the account now?
- Is the bank happy with our implementation?
 - What if this is not withdraw, but deposit?

How Interleaved Can It Get?

- ▣ We'll assume that the only **atomic** operations are reads and writes of words
 - ◆ Some architectures don't even give you that!
- ▣ We'll assume that a **context switch** can occur at any time
- ▣ We'll assume that **you can** delay a thread as long as you like, provided that it's not delayed forever

```
..... get_balance(account);
```

```
balance = get_balance(account);
```

```
balance = .....
```

```
balance = balance - amount;
```

```
balance = balance - amount;
```

```
put_balance(account, balance);
```

```
put_balance(account, balance);
```

Mutual Exclusion

- ▣ We want to use **mutual exclusion** to synchronize access to shared resources
 - ◆ This allows us to have larger atomic code blocks
- ▣ Code that uses mutual exclusion to synchronize its execution is called a **critical region, or critical section**
 - ◆ Only one thread at a time can execute in the critical region
 - ◆ All other threads are forced to wait on entry
 - ◆ When a thread leaves a critical region, another can enter
 - ◆ Example: **sharing your bathroom with roommates**

Critical Region

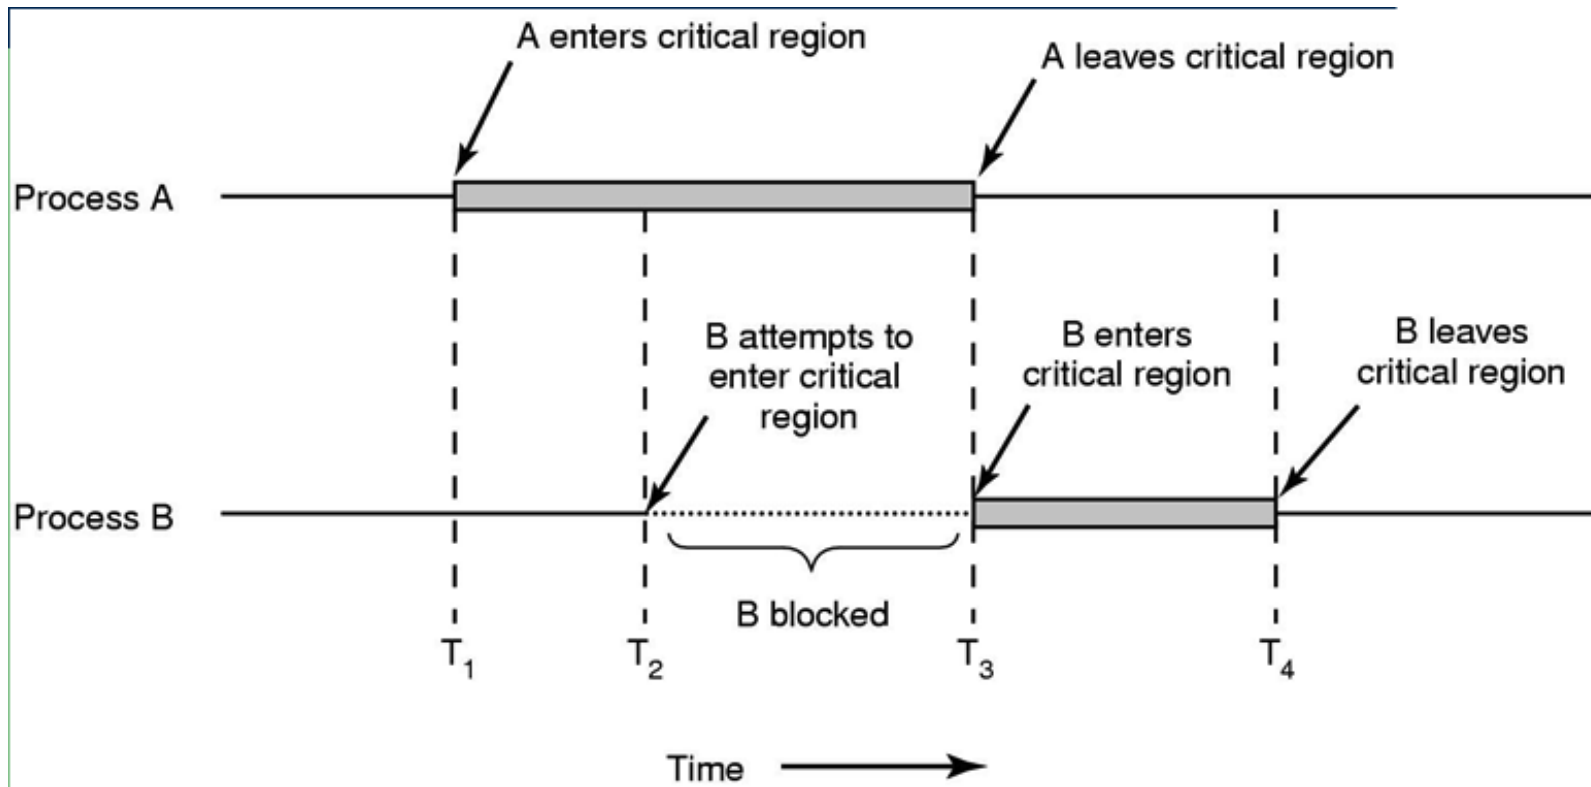
```
Process {  
    while (true) {  
        ENTER CRITICAL SECTION  
        Access shared variables; // Critical Section;  
        LEAVE CRITICAL SECTION  
        Do other work  
    }  
}
```

What requirements would you place on a critical region?

Critical Region Requirements (for thread and process)

- ▣ Mutual exclusion (mutex)
 - ◆ No other thread can execute in the critical region while a thread is in it
- ▣ Progress
 - ◆ A thread in the critical region will eventually leave
 - ◆ If some thread T is not in the critical region, T cannot prevent another thread S from entering the critical region
- ▣ Bounded waiting (no starvation)
 - ◆ If thread T is waiting on the critical region, it will not wait indefinitely
- ▣ No assumption
 - ◆ No assumption may be made about the speed or number of CPUs

Critical Region Illustrated



Mechanisms For Building Critical Regions

- ▣ Atomic read/write
 - ◆ Cannot be “interrupted” when running -> “All or nothing”
- ▣ Locks
 - ◆ Primitive, minimal semantics, used to build others
- ▣ Semaphores
 - ◆ Basic, easy to get the hang of, but hard to program with
- ▣ Monitors
 - ◆ High-level, requires language support, operations implicit
- ▣ Messages
 - ◆ Simple model of communication and synchronization
 - ◆ Direct application to distributed systems

Critical Region with Atomic Read/Write: First Try

```
int turn = 1;
```

```
while (true) {  
    while (turn != 1) ;  
    critical region  
    turn = 2;  
    outside of critical region  
}
```

```
while (true) {  
    while (turn != 2) ;  
    critical region  
    turn = 1;  
    outside of critical region  
}
```

This is called **alternation**

It **satisfies mutex**:

- If blue is in the critical region, then $\text{turn} == 1$ and if yellow is in the critical region, then $\text{turn} == 2$ (why?)
- $(\text{turn} == 1) \equiv (\text{turn} != 2)$

It **violates progress**: the thread could go into an infinite loop outside of the critical section, which will prevent the yellow one from entering.

Easy to use? (what if more than 2 threads? what if we don't know how many threads?)

Locks

- ▣ A lock is an object in memory providing two operations
 - ◆ `acquire()`: before entering the critical region
 - ◆ `release()`: after leaving a critical region
- ▣ Threads **pair** the `acquire()` and `release()` calls
 - ◆ Between `acquire()`/`release()`, the thread **holds** the lock
 - ◆ `acquire()` does not return until any previous holder releases
 - ◆ What can happen if the calls are **not paired**?
- ▣ Locks can spin (a spinlock) or block (a mutex)

Using Locks

```
withdraw (account, amount) {  
    acquire(lock);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    release(lock);  
    return amount;  
}
```

Critical
Region

```
acquire(lock);  
balance = get_balance(account);  
balance = balance - amount;
```

```
acquire(lock);
```

```
put_balance(account, balance);  
release(lock);
```

```
balance = get_balance(account);  
balance = balance - amount;  
put_balance(account, balance);  
release(lock);
```

- What happens when blue tries to acquire the lock?
- Why is the "return" outside the critical region? Is this OK?
- What happens when a third thread calls acquire?

Implementing Locks (1)

- How do we implement locks? Here is one attempt:

```
struct lock {  
    int held = 0;  
}  
  
void acquire (lock) {  
    while (lock->held);  
    lock->held = 1;  
}  
  
void release (lock) {  
    lock->held = 0;  
}
```

busy-wait (spin-wait) for
lock to be released

- This is called a **spinlock** because a thread spins waiting for the lock to be released
 - Does this work?

Implementing Locks (2)

- ❑ No! Two independent threads may both notice that a lock has been released and thereby acquire it.

```
struct lock {  
    int held = 0;  
}  
void acquire (lock) {  
    while (lock->held);  
    lock->held = 1;  
}  
void release (lock) {  
    lock->held = 0;  
}
```

A context switch can occur here,
causing a race condition

Implementing Locks (3)

- ▣ The problem is that the implementation of locks has critical sections, too
 - ◆ How do we stop the recursion?
- ▣ The implementation of acquire/release must be **atomic**
 - ◆ "All or nothing"
- ▣ How do we make them atomic?
- ▣ Need help from hardware
 - ◆ Atomic instructions (e.g., test-and-set)
 - ◆ Disable/enable interrupts (prevents context switches)

Atomic Instructions: Test-And-Set

- The semantics of test-and-set are:

- ◆ Record the old value
- ◆ Set the value to **TRUE**
- ◆ Return the old value

- Hardware executes it **atomically**!

- When executing test-and-set on “flag”

- ◆ What is the **value of flag** afterwards if it was initially False?
- ◆ What is the **return result** if flag was initially False?

```
bool test_and_set (bool *flag) {  
    bool old = *flag;  
    *flag = True;  
    return old;  
}
```


Using Test-And-Set

- ▣ Here is our lock implementation with test-and-set:

```
struct lock {  
    int held = 0;  
}  
void acquire (lock) {  
    while (test-and-set(&lock->held));  
}  
void release (lock) {  
    lock->held = 0;  
}
```

- ▣ When will the while return? What is the value of held?
- ▣ Does it work? What about multiprocessors?

Problems with Spinlocks

- ▣ The problem with spinlocks is that they are wasteful
 - ◆ If a thread is spinning on a lock, then the thread holding the lock cannot make progress
- ▣ Solution 1:
 - ◆ If cannot get the lock, call `thread_yield` to give up the CPU
- ▣ Solution 2: sleep and wakeup
 - ◆ When blocked, go to sleep
 - ◆ Wakeup when it is OK to retry entering the critical region

Disabling Interrupts

- Another implementation of acquire/release is to disable interrupts:

```
struct lock {  
}  
void acquire (lock) {  
    disable interrupts;  
}  
void release (lock) {  
    enable interrupts;  
}
```

- Note that there is no state associated with the lock
- Can two threads disable interrupts simultaneously?

On Disabling Interrupts

- ▣ Disabling interrupts blocks notifications of external events that could trigger a context switch (e.g., timer)
- ▣ In a “real” system, this is only available to the kernel
 - ◆ Why?
- ▣ Disabling interrupts is insufficient on a multiprocessor
 - ◆ Back to atomic instructions

Critical regions without hardware support?

- ❑ So far, we have seen how to implement critical regions (lock) with hardware support
 - ◆ Atomic instruction
 - ◆ Disabling interrupts
- ❑ Can we implement lock *without* HW support?
 - ◆ Software only solution
- ❑ Yes, but...
 - ◆ Complicated (easy to make mistake)
 - ◆ Poor performance
 - ◆ Production OSes use hardware support

Mutex without hardware support: Peterson's Algorithm

```
int turn = 1;  
bool try1 = false, try2 = false;
```

```
while (true) {  
    try1 = true;  
    turn = 2;  
    while (try2 && turn != 1) ;  
    critical section  
    try1 = false;  
    outside of critical section  
}
```

```
while (true) {  
    try2 = true;  
    turn = 1;  
    while (try1 && turn != 2) ;  
    critical section  
    try2 = false;  
    outside of critical section  
}
```

Did I execute "turn=2" before thread 2 executed "turn=1"?

Has thread 2 executed "try2=true?". If not, I am safe. If yes, let's see...


- Does it work?
 - Yes!
- Try all possible interleavings

Summarize Where We Are

- Goal: Use **mutual exclusion** to protect **critical sections** of code that access **shared resources**
- Method: Use locks (spinlocks or disable interrupts)
- Problem: Critical sections can be long

Spinlocks:

- Threads waiting to acquire lock spin in test-and-set loop
- Wastes CPU cycles
- Longer the CS, the longer the spin
- Greater the chance for lock holder to be interrupted



Disabling Interrupts:

- Should not disable interrupts for long periods of time
- Can miss or delay important events (e.g., timer, I/O)

Takeaway from this lecture

- ▣ When you have concurrency and shared resources, protect your critical region with synchronization primitives (e.g., locks, semaphore (next lecture), etc.)