

CSCI3150 Introduction to Operating Systems

Lecture 8: Common Concurrency Problems

Hong Xu

<https://github.com/henryhxu/CSCI3150>

Non-Deadlock Bugs

- ▣ Makes up a majority of concurrency bugs.
- ▣ Two major types of non-deadlock bugs:
 - ◆ Atomicity violation
 - ◆ Order violation

Atomicity-Violation Bugs

- ▣ The desired **serializability** among multiple memory accesses is *violated*.
 - ◆ Simple Example found in MySQL:
 - Two different threads access the field `proc_info` in the struct `thd`.

```
1  Thread1::  
2  if(thd->proc_info){  
3      ...  
4      fputs(thd->proc_info , ...);  
5      ...  
6  }  
7  
8  Thread2::  
9  thd->proc_info = NULL;
```

Atomicity-Violation Bugs (Cont.)

- ▣ **Solution:** Simply add locks around the shared-variable references.

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Thread1::
4  pthread_mutex_lock(&lock);
5  if(thd->proc_info){
6      ...
7      fputs(thd->proc_info , ...);
8      ...
9  }
10 pthread_mutex_unlock(&lock);
11
12 Thread2::
13 pthread_mutex_lock(&lock);
14 thd->proc_info = NULL;
15 pthread_mutex_unlock(&lock);
```

Order-Violation Bugs

- ▣ The **desired order** between two memory accesses is flipped.
 - ◆ i.e., **A** should always be executed before **B**, but the order is not enforced during execution.
 - ◆ **Example:**
 - The code in Thread2 seems to assume that the variable `mThread` has already been *initialized* (and is not `NULL`).

```
1  Thread1::  
2  void init() {  
3      mThread = PR_CreateThread(mMain, ...);  
4  }  
5  
6  Thread2::  
7  void mMain(...) {  
8      mState = mThread->State  
9  }
```

Order-Violation Bugs (Cont.)

- ▣ **Solution:** Enforce ordering using **condition variables**

```
1  pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;
2  pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;
3  int mtInit = 0;
4
5  Thread 1::
6  void init(){
7      ...
8      mThread = PR_CreateThread(mMain,...);
9
10     // signal that the thread has been created.
11     pthread_mutex_lock(&mtLock);
12     mtInit = 1;
13     pthread_cond_signal(&mtCond);
14     pthread_mutex_unlock(&mtLock);
15     ...
16 }
17
```

Order-Violation Bugs (Cont.)

```
18 Thread2::  
19 void mMain(...) {  
21     // wait for the thread to be initialized ...  
22     pthread_mutex_lock(&mtLock);  
23     while(mtInit == 0)  
24         pthread_cond_wait(&mtCond, &mtLock);  
25     pthread_mutex_unlock(&mtLock);  
26  
27     mState = mThread->State;  
28     ...  
29 }
```

Deadlock Bugs

Thread 1:

```
lock(L1);
```

```
lock(L2);
```


Thread 2:

```
lock(L2);
```

```
lock(L1);
```

- ◆ The presence of a cycle

- Thread1 is holding a lock L1 and waiting for another one, L2.
- Thread2 that holds lock L2 is waiting for L1 to be released.



Why Do Deadlocks Occur?

▣ Reason 1:

- ◆ In large code bases, **complex dependencies** arise between components.

▣ Reason 2:

- ◆ Due to the nature of **encapsulation**
 - Hide details of implementations and make software easier to build in a modular way.
 - Such **modularity** *does not mesh* well with locking.

Why Do Deadlocks Occur? (Cont.)

- ▣ **Example:** Java Vector class and the method `AddAll()`

```
1  vector v1,v2;  
2  // thread 1  
3  v1.AddAll(v2);  
  
4  // thread 2  
5  v2.AddAll (v1) ;
```

- ◆ **Locks** for both the vectors being added to (`v1`) and the parameter (`v2`) *need to be acquired*.

Conditions for Deadlock

- ▣ Four conditions need to hold for a deadlock to occur.

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them.
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

- ◆ If any of these four conditions are not met, **deadlock cannot occur**.

Prevention – Circular Wait

- ▣ Provide a total ordering on lock acquisition
 - ◆ This approach requires *careful design* of global locking strategies.
- ▣ **Example:**
 - ◆ There are two locks in the system (L1 and L2)
 - ◆ We can prevent deadlock by always acquiring L1 before L2.

Avoiding Circular Wait

- <https://elixir.bootlin.com/linux/latest/source/mm/filemap.c>

```
/*
 * Shared mappings implemented 30.11.1994. It's not fully working yet,
 * though.
 *
 * Shared mappings now work. 15.8.1995 Bruno.
 *
 * finished 'unifying' the page and buffer cache and SMP-threaded the
 * page-cache, 21.05.1999, Ingo Molnar <mingo@redhat.com>
 *
 * SMP-threaded pagemap-LRU 1999, Andrea Arcangeli <andrea@suse.de>
 */

/*
 * Lock ordering:
 *
 * ->i_mmap_rwsem          (truncate_pagecache)
 *   ->private_lock        (__free_pte->__set_page_dirty_buffers)
 *     ->swap_lock          (exclusive_swap_page, others)
 *       ->i_pages lock
 *
 * ->i_mutex
 *   ->i_mmap_rwsem          (truncate->unmap_mapping_range)
 *
 * ->mmap_sem
 *   ->i_mmap_rwsem
 *     ->page_table_lock or pte_lock  (various, mainly in memory.c)
 *       ->i_pages lock          (arch-dependent flush_dcache_mmap_lock)
```

Prevention – Hold-and-wait

- ▣ Acquire all locks **at once, atomically**.

```
1    lock(prevention) ;  
2    lock(L1) ;  
3    lock(L2) ;  
4    ...  
5    unlock(prevention) ;
```

- ◆ This code guarantees that **no untimely thread switch can occur** *in the midst of* lock acquisition.
- ◆ **Problem:**
 - Require us to know when calling a routine exactly which locks must be held and to acquire them ahead of time.
 - Decrease *concurrency*

Prevention – No Preemption

- ❑ **Multiple lock acquisition** often gets us into trouble because when waiting for one lock **we are holding another**.
- ❑ `trylock()`
 - ◆ Used to build a *deadlock-free, ordering-robust* lock acquisition protocol.
 - ◆ Grab the lock (if it is available).
 - ◆ Or, return -1: you should try again later.

```
1  top:
2      lock(L1);
3      if( tryLock(L2) == -1 ){
4          unlock(L1);
5          goto top;
6      }
```

Prevention – No Preemption (Cont.)

▣ livelock

- ◆ Both systems are running through the code sequence *over and over again*.
- ◆ Progress is not being made.
- ◆ Solution:
 - Add a **random delay** before looping back and trying the entire thing over again.

Prevention – Mutual Exclusion

□ wait-free

- ◆ Using powerful **hardware instructions (CPU)**.
- ◆ You can build data structures in a manner that *does not require explicit locking*.

```
1  int CompareAndSwap(int *address, int expected, int new){
2      if(*address == expected){
3          *address = new;
4          return 1; // success
5      }
6      return 0;
7  }
```

Prevention – Mutual Exclusion (Cont.)

- We now want to **atomically increment** a value by a certain amount:

```
1  void AtomicIncrement(int *value, int amount) {  
2      do {  
3          int old = *value;  
4      } while ( CompareAndSwap(value, old, old+amount) == 0 );  
5  }
```

- ◆ Repeatedly tries to update the value to *the new amount* and uses the compare-and-swap to do so.
- ◆ **No lock** is acquired
- ◆ **No deadlock** can arise
- ◆ **livelock** is still a possibility.

Prevention – Mutual Exclusion (Cont.)

▣ Another example:


- ◆ Surrounding this code with a **lock acquire** and **release**.

```
1  void insert(int value){
2      node_t * n = malloc(sizeof(node_t));
3      assert( n != NULL );
4      n->value = value ;
5      lock(listlock); // begin critical section
6      n->next = head;
7      head     = n;
8      unlock(listlock) ; //end critical section
9  }
```


- ◆ **wait-free manner** using the compare-and-swap instruction

```
1  void insert(int value) {
2      node_t *n = malloc(sizeof(node_t));
3      assert(n != NULL);
4      n->value = value;
5      do {
6          n->next = head;
7      } while (CompareAndSwap(&head, n->next, n));
8  }
```

Insert with lock



Lock-free insert



Deadlock Avoidance via Scheduling

- ❑ Deadlock Avoidance
 - ◆ Get the information about the locks various threads might grab during their execution.
 - ◆ schedule the threads in a way to guarantee no deadlock can occur.
- ❑ In some scenarios, **deadlock avoidance** is preferable.
- ❑ Problem: Global knowledge is required.
- ❑ Not commonly used. Refer to textbook for more details

Detect and Recover

- **Allow deadlocks** to occasionally occur and then *take some action*.
 - ◆ **Example:** if an OS froze, you would reboot it.
- Many database systems employ *deadlock detection* and *recovery technique*.
 - ◆ A deadlock detector **runs periodically**.
 - ◆ Building a **resource graph** and checking it for cycles.
 - ◆ In deadlock, the system **needs to be restarted**.

Summary

- ▣ Non-deadlocked bugs
 - ◆ Atomicity violation
 - ◆ Ordering violation
- ▣ Deadlock
 - ◆ Enforcing the lock order
 - ◆ Lockless mechanism