

CSCI3150 Introduction to Operating Systems

Lecture 10: Memory Management I

Hong Xu

<https://github.com/henryhxu/CSCI3150>

Overview

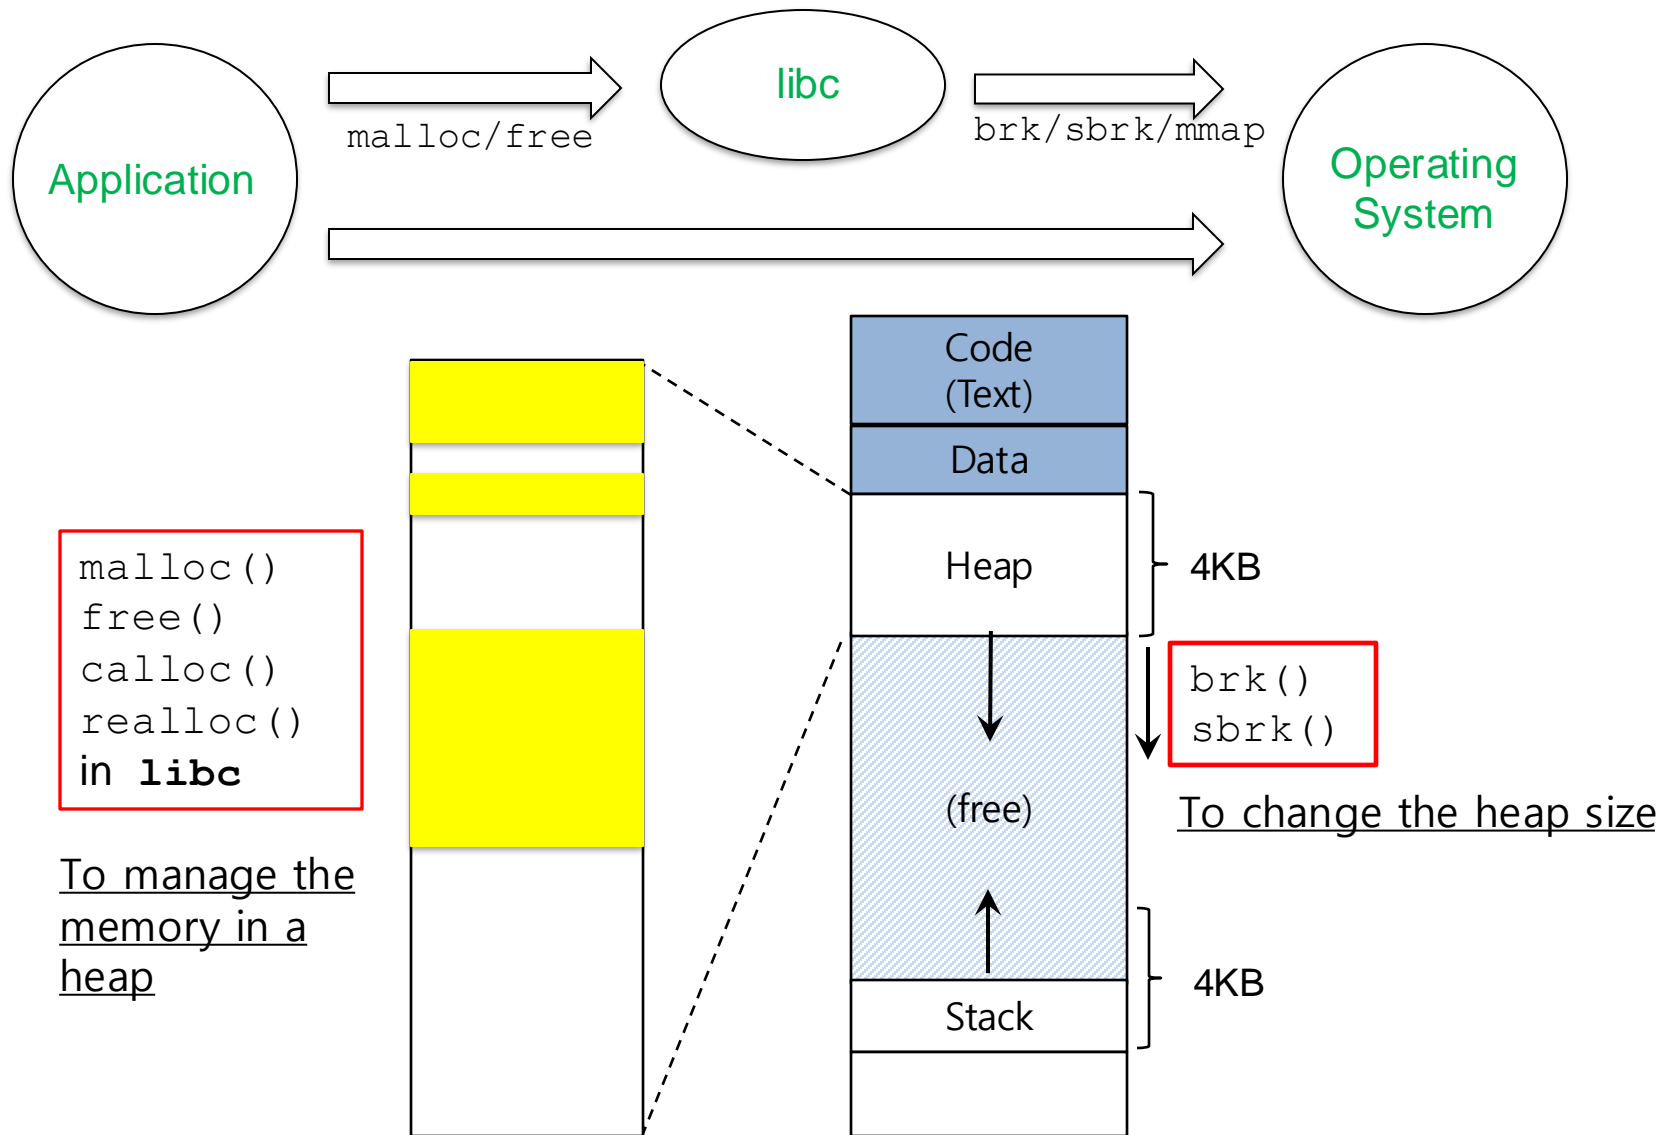
- ▣ Address space, memory API
- ▣ Address translation: Base-and-bounds
- ▣ Segmentation

Memory API

Overview

- ▣ `malloc/free`
- ▣ `calloc/realloc`
- ▣ `brk/sbrk`

Virtual Address Space



malloc()

```
#include <stdlib.h>

void* malloc(size_t size)
```

- ▣ Allocate a memory region on the heap.
 - ◆ Argument
 - `size_t size`: size of the memory block (in bytes)
 - `size_t` is an unsigned integer type
 - ◆ Return
 - Success: a void type pointer to the memory block allocated by `malloc`
 - Fail: a null pointer

sizeof()

- ▣ Routines and macros are utilized for `size` in `malloc` instead typing in a number directly.
- ▣ Two types of results when using `sizeof` on variables; **be careful...**
 - ◆ The actual size of `'x'` is known at run-time.

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

4

- ◆ The actual size of `'x'` is known at compile-time.

```
int x[10];  
printf("%d\n", sizeof(x));
```

40


Memory API: `free()`

```
#include <stdlib.h>

void free(void* ptr)
```


- ▣ Free a memory region allocated by a call to `malloc`
 - ◆ Argument
 - `void *ptr`: a pointer to a memory block allocated with `malloc`
 - ◆ Return
 - none

Memory Allocating




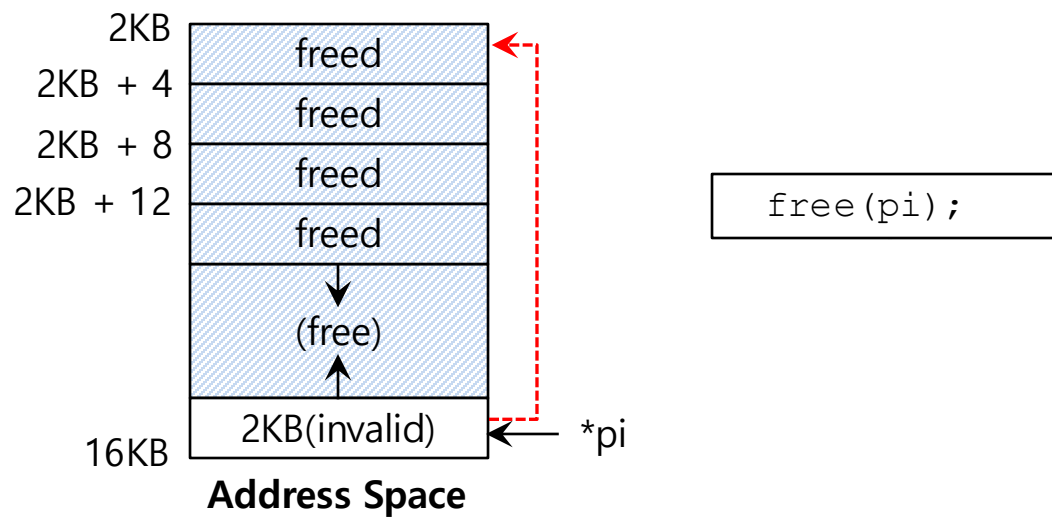
-----> pointer

```
int *pi; // local variable
```



```
pi = (int *) malloc(sizeof(int) * 4);
```

Memory Freeing

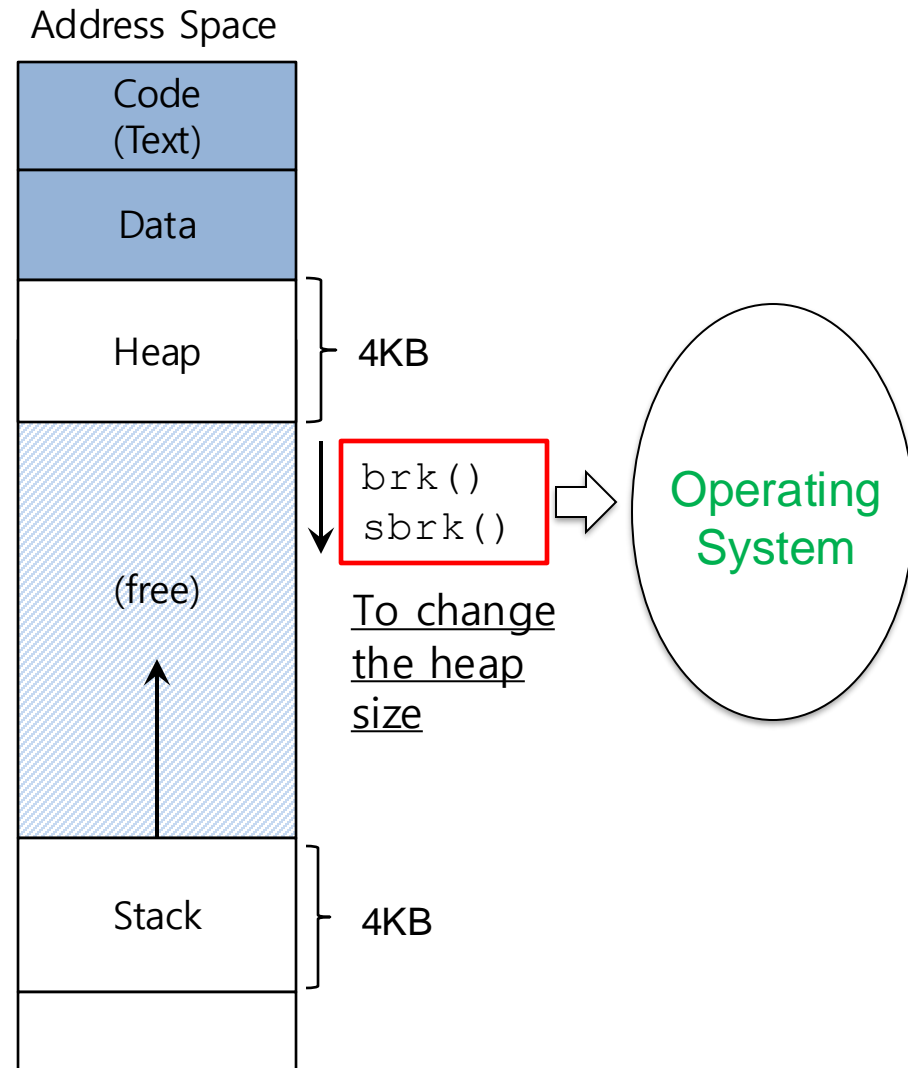


System Calls

```
#include <unistd.h>

int brk(void *addr)
void *sbrk(intptr_t increment);
```

- ▣ There lacks heap space → Ask OS to expand the heap
- ▣ `break`: The location of **the end of the heap** in address space
- ▣ `malloc` uses `brk` syscall
 - ◆ `brk` is called to expand the program's *break*.
 - ◆ `sbrk` is similar to `brk`.
 - ◆ Programmers **should never directly call** either `brk` or `sbrk`.



Address Translation

Memory Virtualization with Efficiency and Control

- ▣ Memory virtualization takes a similar strategy known as **limited direct execution (LDE)**
- ▣ Efficiency and control are attained by **hardware support**.
 - ◆ e.g., registers, TLBs (Translation Look-aside Buffers), page table

Address Translation

- ▣ Hardware transforms a **virtual address** to a **physical address**.
 - ◆ The desired information is stored in a physical address.
- ▣ The OS must get involved at key points to set up the hardware.
 - ◆ The OS must manage memory to judiciously intervene.

Example: Address Translation

▣ Code in C

```
void func()  
    int x=3000;  
    ...  
    x = x + 3; // this is the line of code we are interested in
```

- ◆ **Load** a value from memory
- ◆ **Increment** it by three
- ◆ **Store** the value back into memory


Example: Address Translation (Cont.)

▣ Assembly

```
128 : movl 0x0(%ebx), %eax      ; load 0+ebx into eax
132 : addl $0x03, %eax         ; add 3 to eax register
135 : movl %eax, 0x0(%ebx)      ; store eax back to mem
```

- ◆ Presume that the address of 'x' has been placed in `ebx` register
- ◆ **Load** the value at that address into `eax`
- ◆ **Add** 3 to `eax`
- ◆ **Store** the value in `eax` back into memory

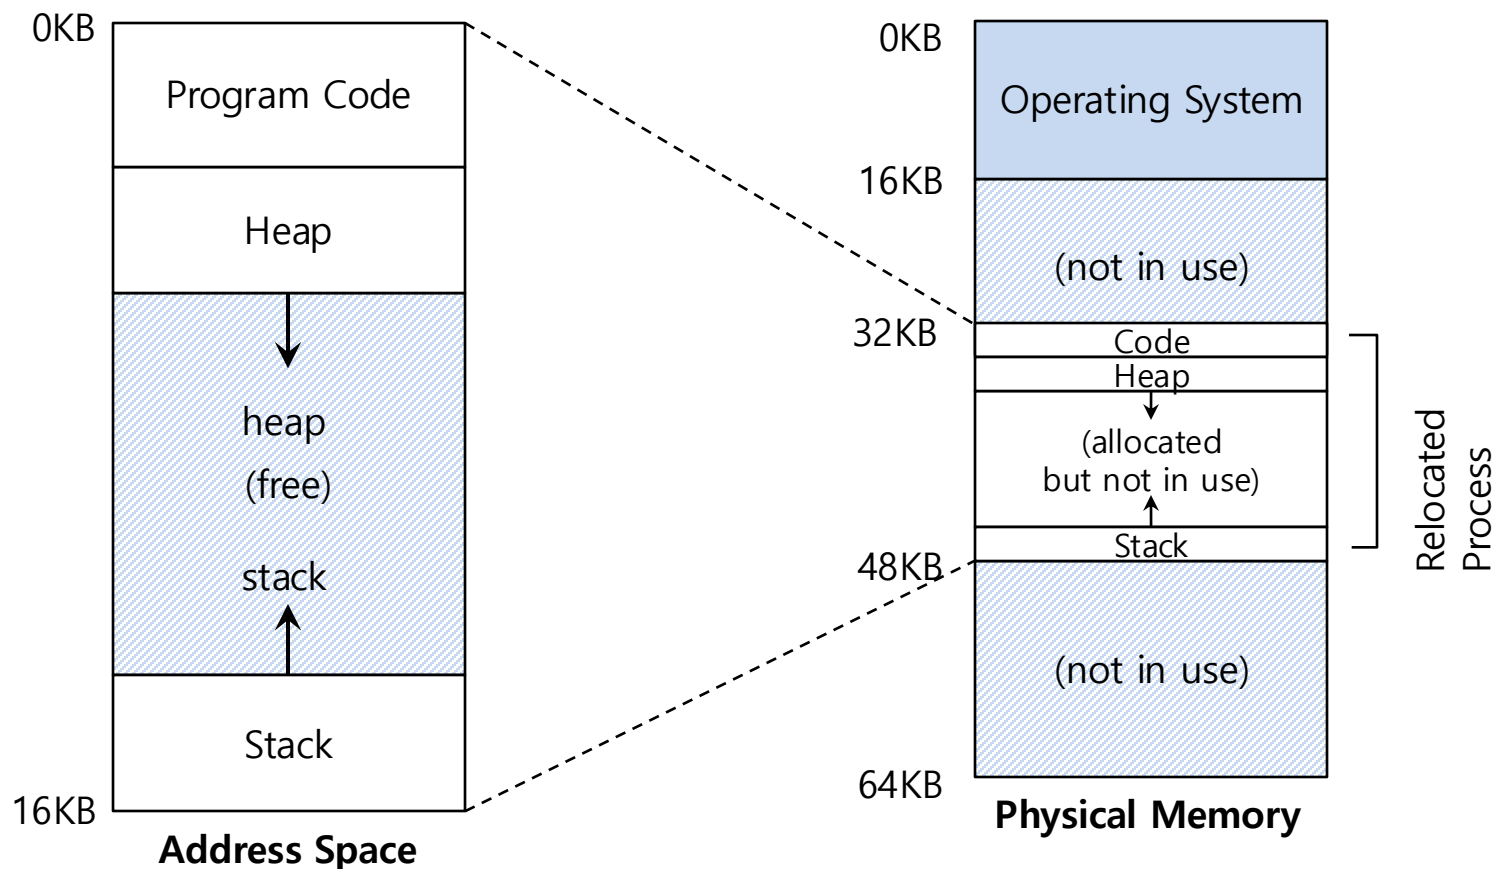
Example: Address Translation (Cont.)



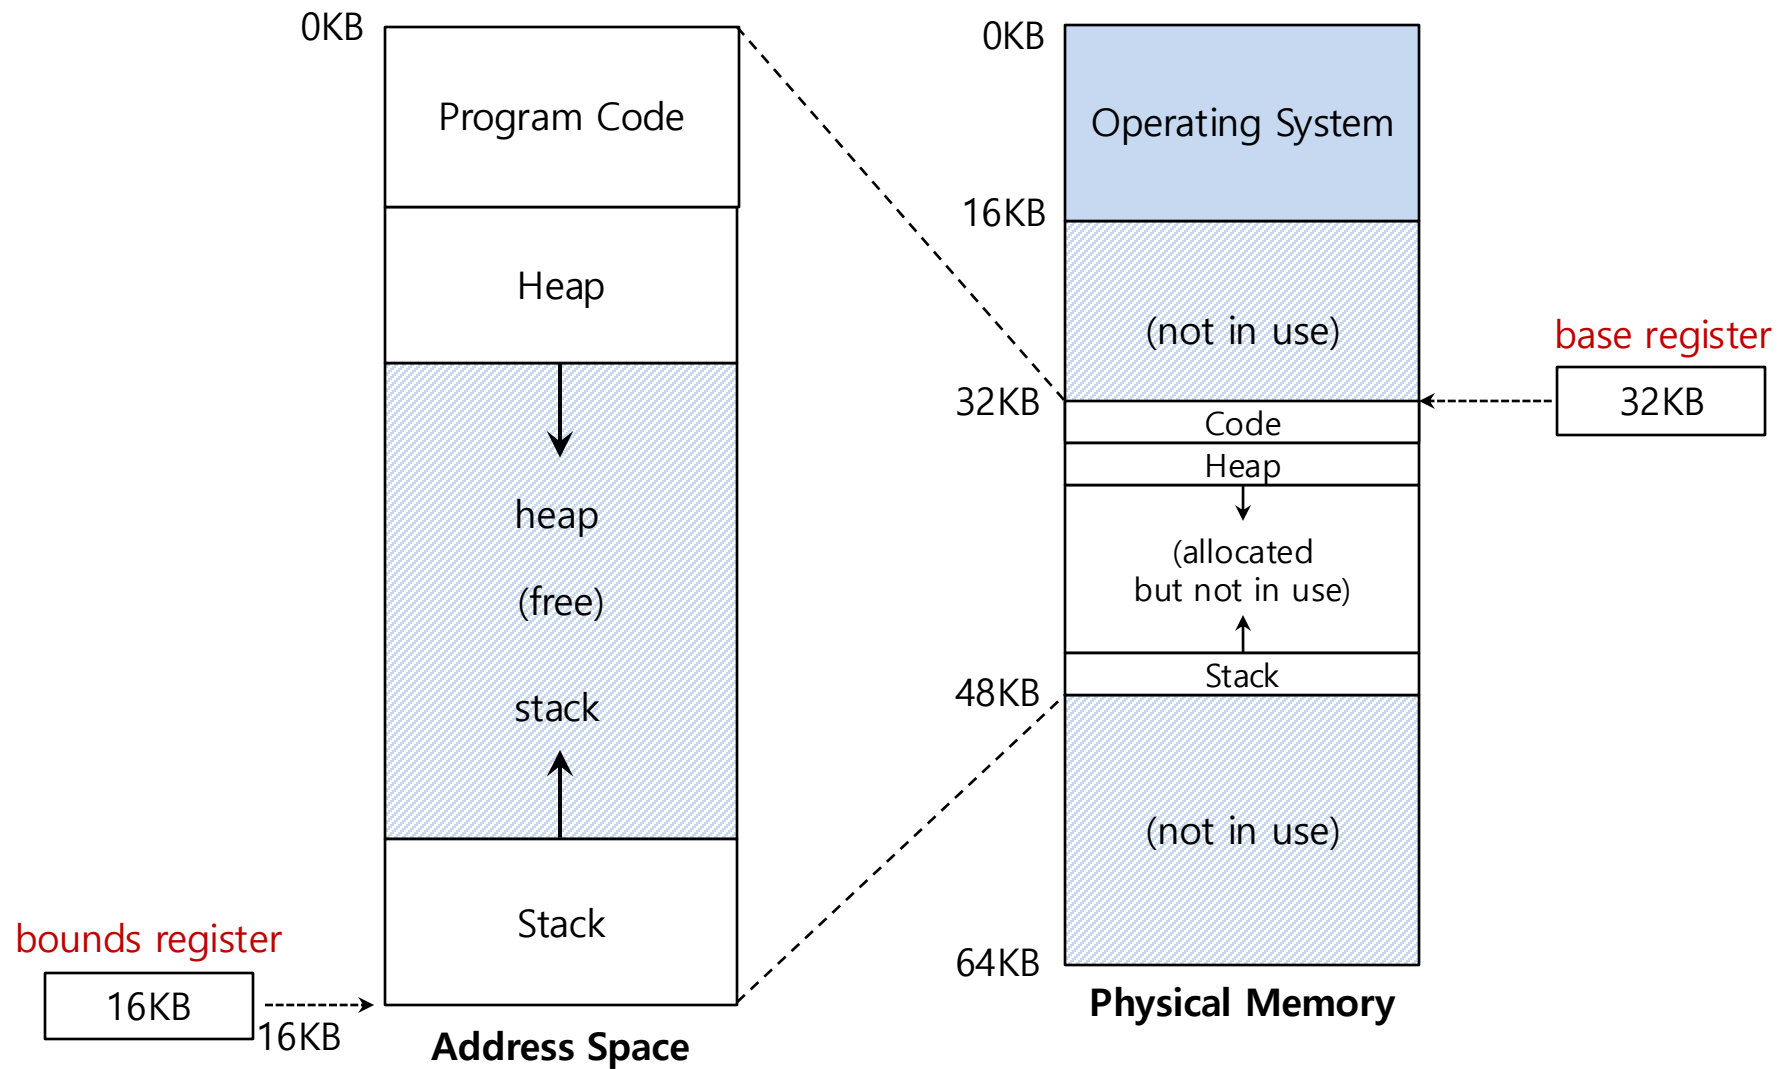
- Fetch instruction at address 128
- Execute this instruction (load from address 15KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)

Dynamic Relocation (Hardware based): Base-and-Bounds

- The OS wants to place the process **somewhere** in the physical memory, not at address 0.
 - ◆ The address space start at address 0.



Base and Bounds Register



Base and Bounds

- When a program starts running, the OS decides **where** in physical memory a process should be **loaded**
 - ◆ Set the **base** register a value.

$$\text{physical address} = \text{virtual address} + \text{base}$$

- ◆ Every virtual address must **not be greater than bounds, or negative**

$$0 \leq \text{virtual address} < \text{bounds}$$

Relocation and Address Translation

128 : `movl 0x0(%ebx), %eax`

- ◆ **Fetch** instruction at address 128


$$32896 = 128 + 32KB(base)$$

- ◆ **Execute** this instruction

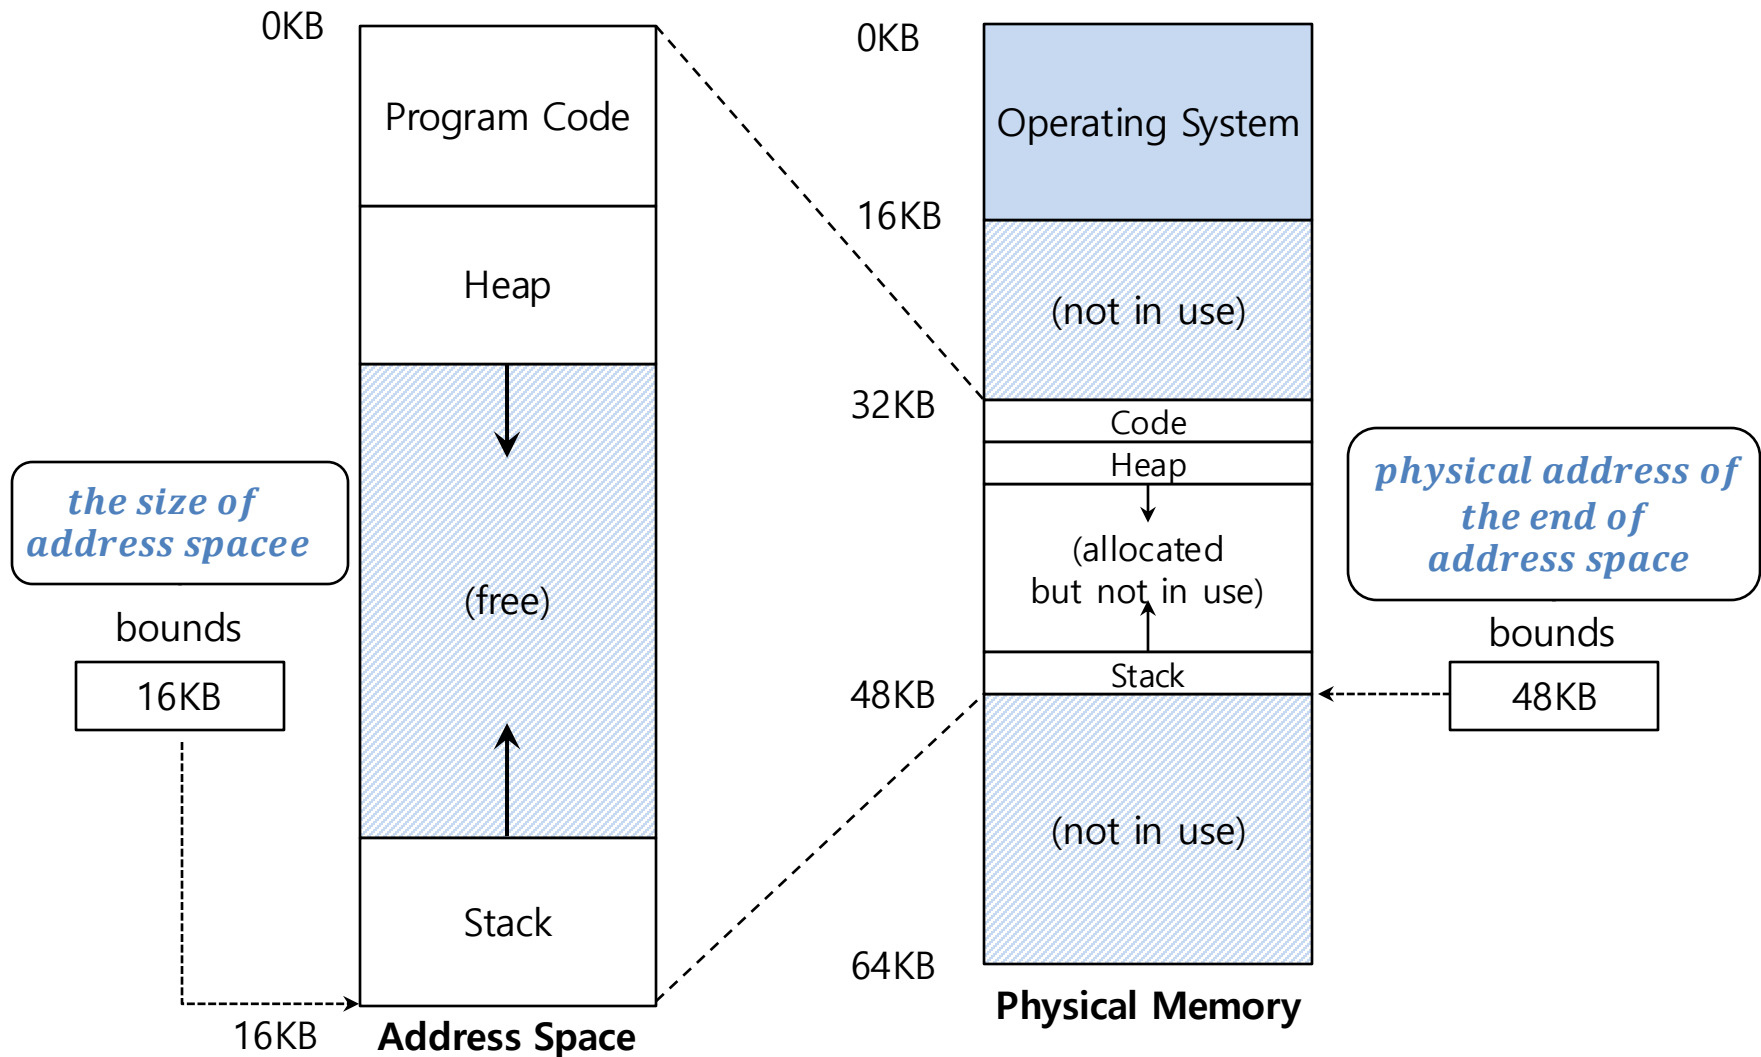
1KB = 1024B

- Load from address 15KB

$$47KB = 15KB + 32KB(base)$$



Two Ways of Bounds Register



Hardware Requirements

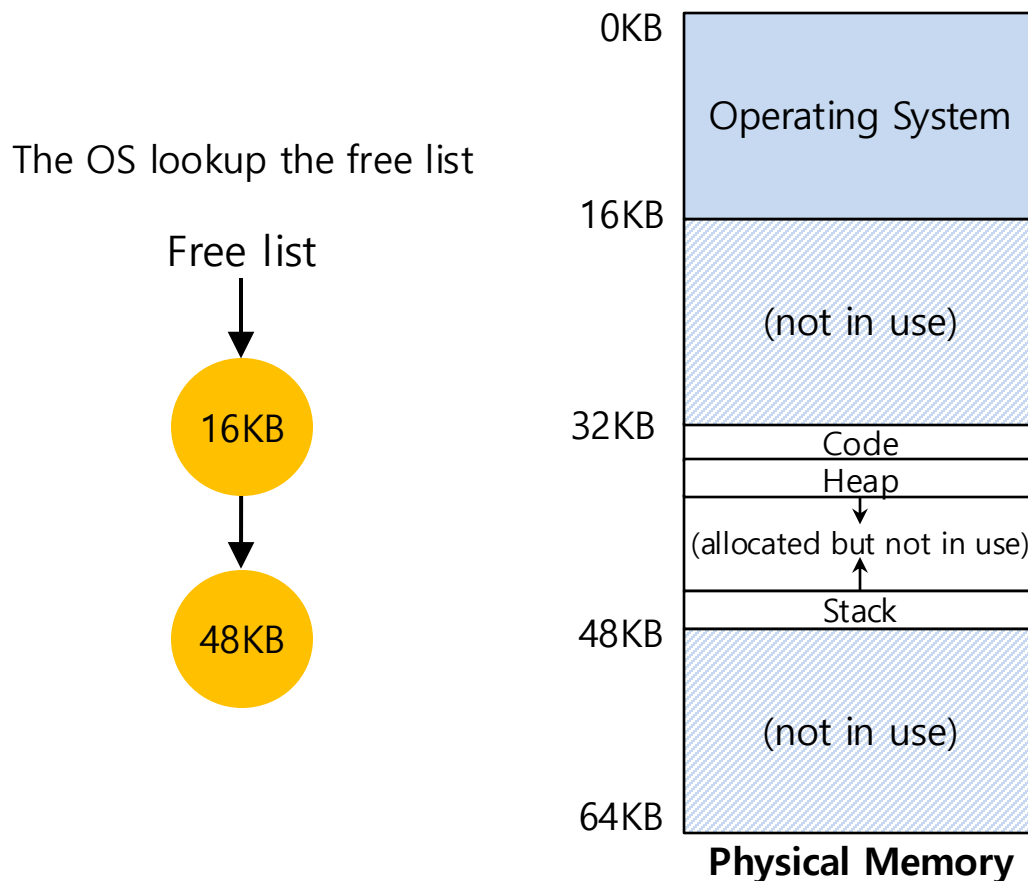
- ▣ Privileged mode: prevent user-mode processes from executing privileged operations
- ▣ Base/bounds registers: Need a pair of registers per CPU to support address translation and bounds checks
- ▣ Ability to translate virtual addresses and check if within bounds limits; Circuitry to do translations.
- ▣ Privileged instruction(s) to update base/bounds: OS must be able to set these values before letting a user program run
- ▣ Privileged instruction(s) to register: OS must be able to tell hardware what exception handlers code to run if exception occurs
- ▣ Ability to raise exceptions when processes try to access privileged instructions or out-of-bounds memory

OS Issues for Memory Virtualization

- ▣ The OS must **take action** to implement **base-and-bounds** approach.
- ▣ Three critical junctures:
 - ◆ When a process **starts running**:
 - Finding space for address space in physical memory
 - ◆ When a process is **terminated**:
 - Reclaiming the memory for use
 - ◆ When context **switch occurs**:
 - Saving and storing the base-and-bounds pair



OS Issues: When a Process Starts Running

- ▣ The OS must **find a room** for a new address space
 - ◆ free list : A list of the ranges of physical memory which are not in use.



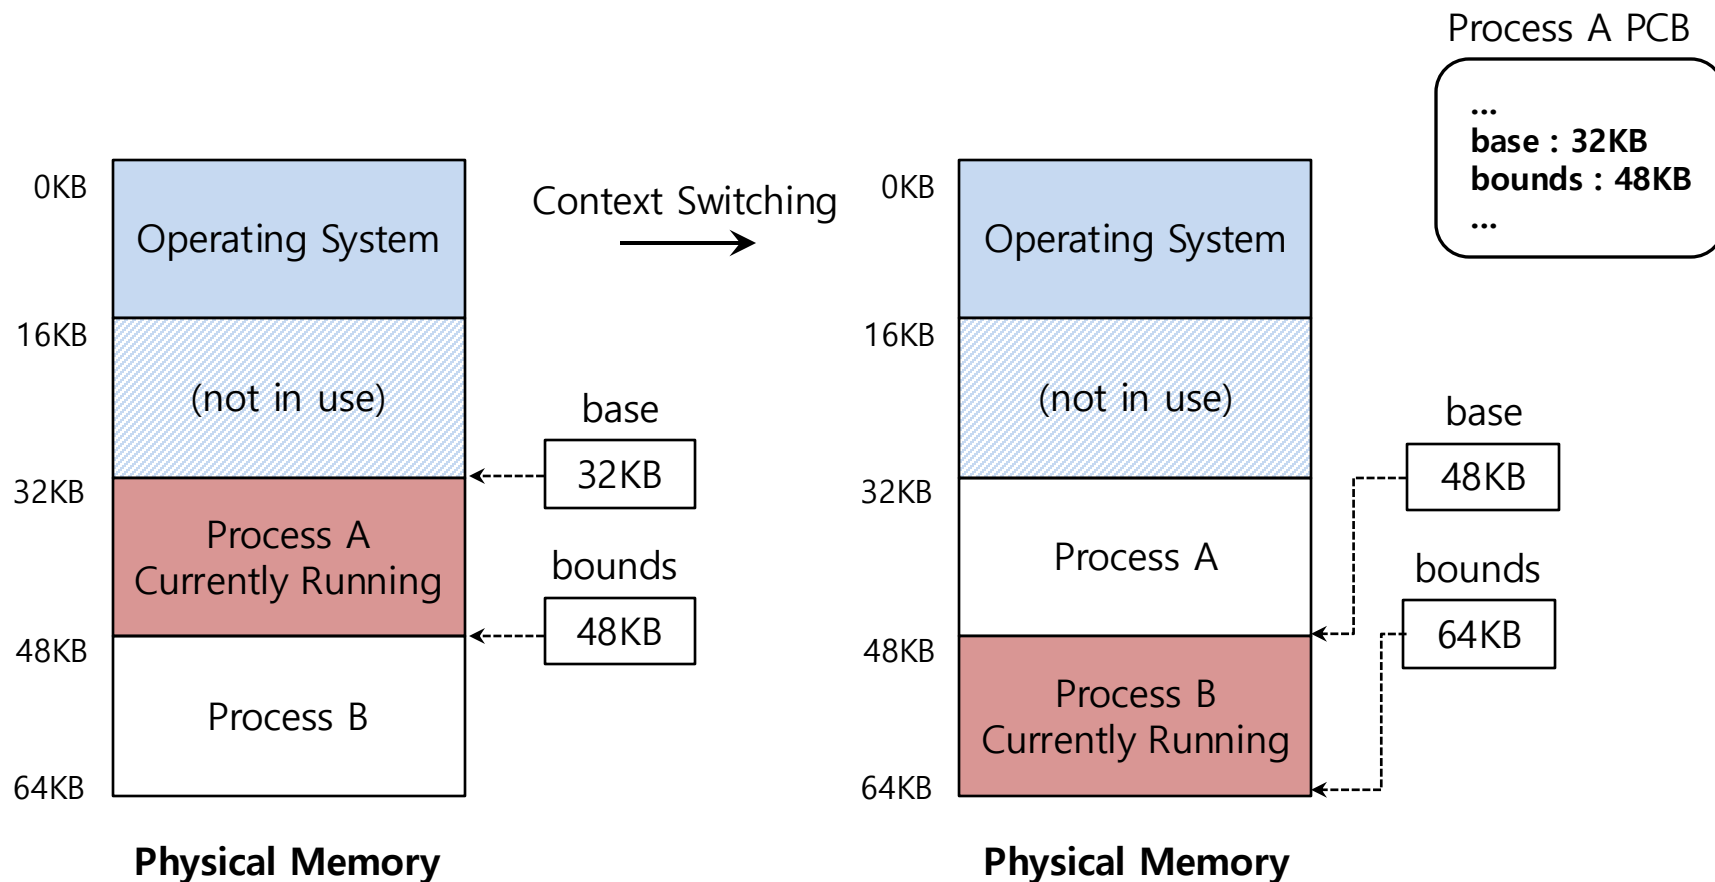
OS Issues: When a Process Is Terminated

- ▣ The OS must **put the memory back** on the free list.




OS Issues: When Context Switch Occurs

- The OS must **save and restore** the base-and-bounds pair.
 - ◆ In **process structure** or **process control block (PCB)**



Segmentation

Inefficiency of Base-and-Bounds




- **Big chunk of “free” space**
- “free” space **takes up** physical memory
- Hard to run when an address space **does not fit** into physical memory

Segmentation

- A **segment** is just a **contiguous portion** of the address space of a particular length
 - ◆ Logically different segments: code, stack, heap
- Each segment can be **placed** in **different parts of the physical memory**
 - ◆ **Base** and **bounds** exist for **each** segment

Placing Segments in Physical Memory

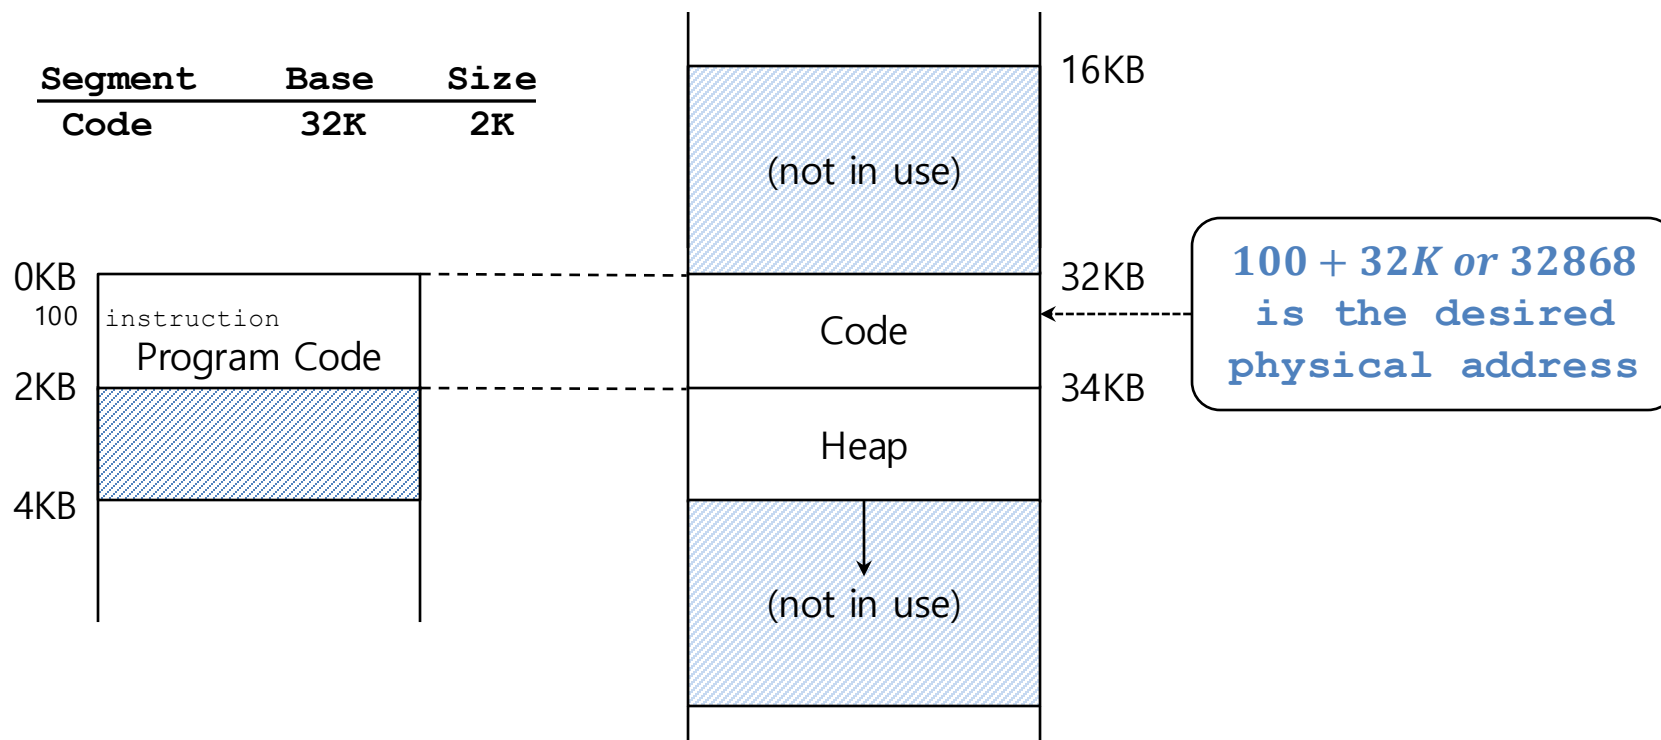


Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Address Translation with Segmentation: code

$$\text{physical address} = \text{offset} + \text{base}$$

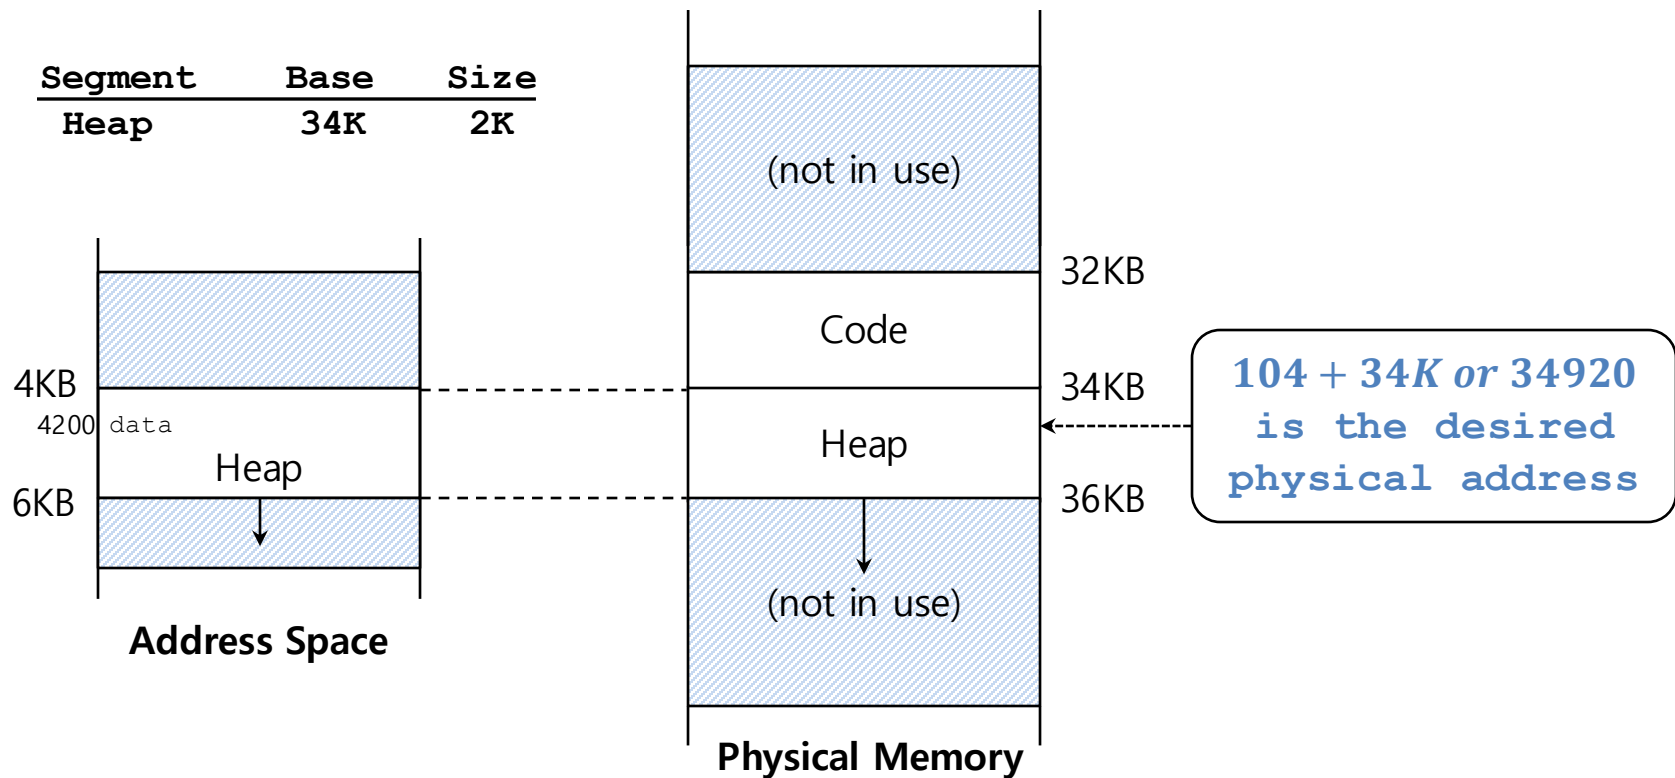
- ▣ The `offset` of virtual address 100 is 100.
 - ◆ The code segment **starts at virtual address 0** in address space.



Address Translation on Segmentation: heap


*Virtual address + base is not the correct physical address.
(Offset of virtual address) + base is the correct physical address.*

- The offset of virtual address 4200 is 104 (4200 - 4096)
 - ◆ The heap segment **starts at virtual address 4096** in address space.



Segmentation Fault or Violation


- If an **illegal address** such as 7KB which is beyond the end of heap is referenced, the OS occurs **segmentation fault**.
 - ◆ The hardware detects that address is **out of bounds**.



Referring to A Segment


□ Explicit approach

- ◆ Chop up the address space into segments based on the **top few bits** of virtual address



□ Example: virtual address 4200 (01000001101000)

Segment	bits
Code	00
Heap	01
Stack	10
-	11




Segment Selection

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

- ◆ `SEG_MASK = 0x3000 (11000000000000)`
- ◆ `SEG_SHIFT = 12`
- ◆ `OFFSET_MASK = 0xFFF (00111111111111)`

Referring to the Stack Segment

- ❑ Stack grows **backwards**
- ❑ **Extra hardware support** is need
 - ◆ The hardware checks which way the segment grows.
 - ◆ 1: positive direction, 0: negative direction



Segment Registers (with Negative-Growth Support)

Segment	Base	Size	Grows Positive?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Support for Sharing

- ▣ Segment can be **shared** between address spaces
 - ◆ **Code sharing** is still in use in systems today
- ▣ Extra hardware support is need for form of **Protection bits**.
 - ◆ **A few more bits** per segment to indicate permissions of read, write and execute

Segment Register Values (with Protection)



Segment	Base	Size	Grows Positive?	Protection
Code	32K	2K	1	Read-Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write

OS Support: Fragmentation

- ▣ **External Fragmentation:** little holes of **free space** in physical memory that may be individually too small for segment
 - ◆ There is **24KB free**, but **not in one contiguous** segment.
 - ◆ The OS **cannot** satisfy a **20KB request**

- ▣ **Compaction: re-arranging** the exiting segments
 - ◆ Compaction is **costly**
 - **Stop** running processes
 - **Copy** data to somewhere
 - **Change** segment register values

Memory Compaction



Issues of Segmentation

- ▣ External fragmentation
- ▣ Still inefficient and inflexible, when a segment is only sparsely used
- ▣ Better solution? Yes!