

# **Assignment One: System Calls in C Programming**

## **CSCI3150 Introduction to Operating Systems, Fall 2024**

Shaofeng Wu  
wsf123@link.cuhk.edu.hk  
Sep 19, 2024

# Overview (1/1)

## ➤ Tutorial last week

- Behavior of a bash shell
  - Built-in commands
  - External commands
  - Pipe
- Code walk of asg1

**All examples used in this tutorial:**  
**<https://github.com/henryhxu/CSCI3150/tree/2024-Fall/tutorial/T03>**

## ➤ Tutorial this week

- System calls in C programming
  - fork & exec
  - dup
  - pipe

# Fork & Exec (1/3)

example1.c

```
#include <sys/wait.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(){

    pid_t pid = 0;
    int status;
    int wait_return;

    for(int i=0; i<10; i++){
        /* create 10 child processes in parent process */
        if(pid == 0){
            pid = fork();
            if(pid < 0){
                printf("fork() failed\n");
                exit(-1);
            }
        }
    }

    if(pid > 0){
        printf("Hello from child %d\n",pid);
        exit(0);
    }

    /* wait for all child to terminate by checking the return value of wait */
    while((wait_return = wait(&status)) > 0){;}

    printf("All child terminated\n");

    return 0;
}
```

## ➤ Fork & wait

### ➤ Function prototype

➤ pid\_t fork(void);

➤ pid\_t wait(int \*wstatus);

### ➤ fork() is used to create child processes

### ➤ wait() is used to let the parent wait for one of its child to terminate

### ➤ The child process created with fork() is an exact duplicate of the parent process except for several points(process id, memory lock, ...)

# Fork & Exec (2/3)

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(){

    char *file_name = "./my_env";
    char *argv[32] = {"my_env", NULL};
    char *envp[2] = { "ABC=1", NULL};

    /* execute the command */

    // execl("/bin/my_env", argv[0], NULL);
    // execlp(file_name, argv[0], NULL);
    // execl("/bin/my_env", argv[0], NULL, envp);

    // execv("/bin/my_env", argv);
    // execvp(file_name, argv);
    execvpe(file_name, argv, envp);

    printf("We just arrive at a line of code after
exec()\n");

    return 0;
}
```

example2.c

## ➤ The exec() system call family:

- int execl(const char \*pathname, const char \*arg, ... /\*, (char \*) NULL \*/);
- int execlp(const char \*file, const char \*arg, ... /\*, (char \*) NULL \*/);
- int execl("/bin/my\_env", const char \*arg, ... /\*, (char \*) NULL, char \*const envp[] \*/);
- int execv(const char \*pathname, char \*const argv[]);
- int execvp(const char \*file, char \*const argv[]);
- int execvpe(const char \*file, char \*const argv[], char \*const envp[]);

# Fork & Exec (3/3)

example3.c

```
#include <sys/wait.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(){

    pid_t pid;
    pid = fork();
    if(pid == -1){
        printf("Fork error\n");
        exit(-1);
    }
    else if(pid == 0){
        /* parent */
        printf("parent is doing something\n");
        sleep(5);
        printf("parent finished its task");
    }
    else{
        /* child */
        execlp("ls","ls","-a",NULL);
    }
    return 0;
}
```

## ➤ Combining fork() and exec()?

- Child process will have its own image, no longer a duplicate of its parent

# Dup (I/I)

```
example4.c
#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>
#include <fcntl.h>

int main(){

    char *file_path = "./test0.txt";
    int fd = open(file_path, O_RDWR);

    /* redirect standard output to the file */
    dup2(fd,STDOUT_FILENO);

    printf("-----dup test ends-----\n");

    close(fd);
    return 0;
}
```

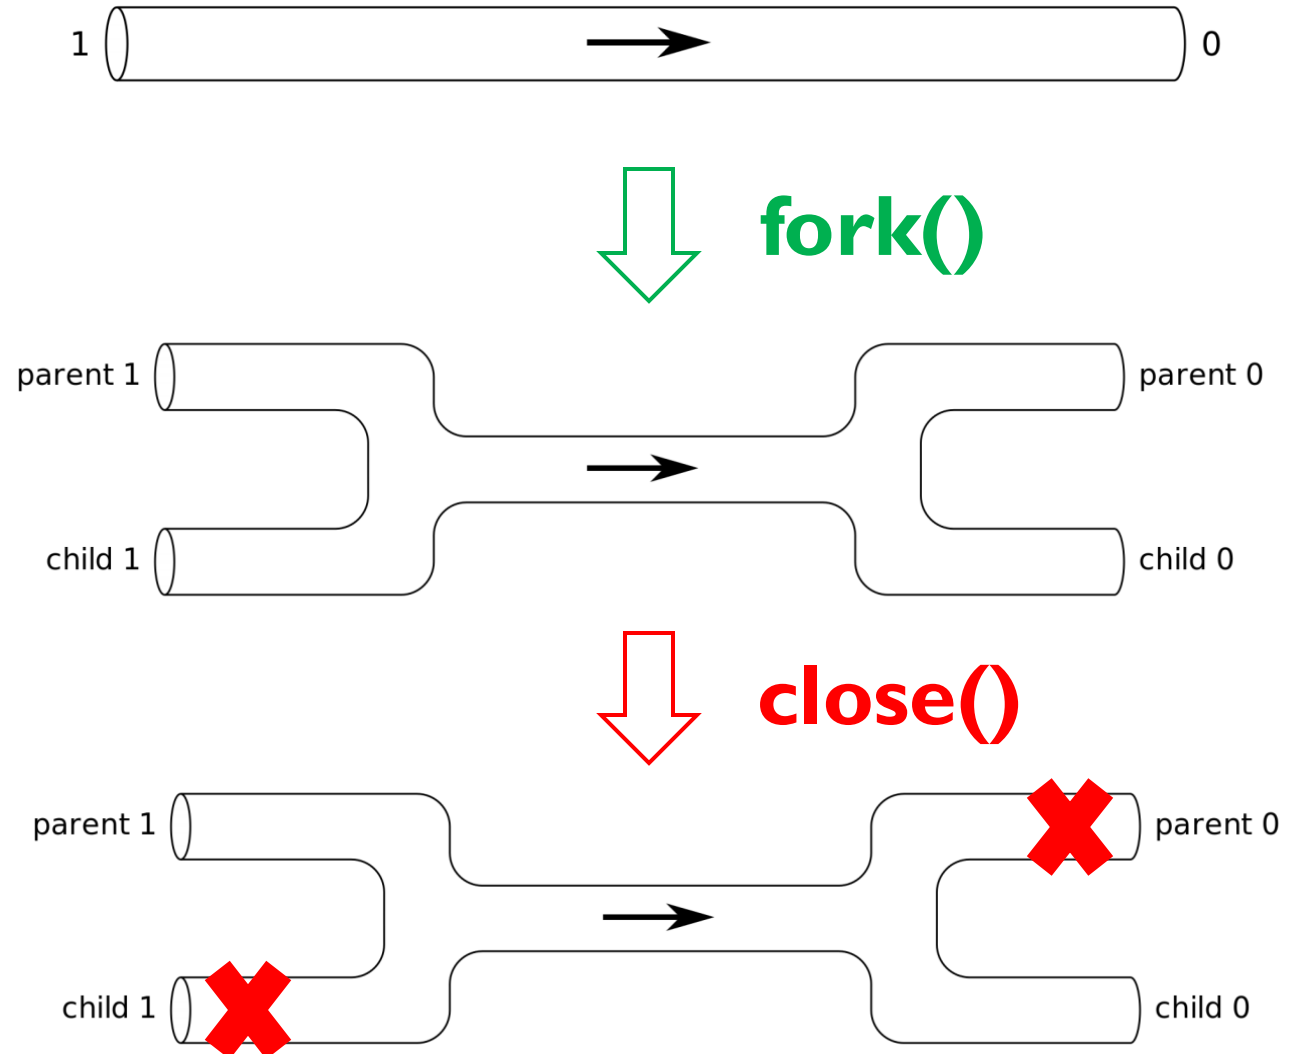
## ➤ dup

- Function prototype
  - int dup(int oldfd);
  - int dup2(int oldfd, int newfd);
- Allocates a new file descriptor that refers to the same open file description as the descriptor oldfd. For dup2, it specifically uses newfd instead of allocating a new file descriptor.
- On success, these system calls return the new file descriptor. On error, -1 is returned.
- Question: How to restore standard output after the dup2()?

# Pipe (1/4)

## ➤ Pipe

- System call for pipe creation
  - `int pipe(int pipefd[2])`
- Pipe characteristics
  - Unidirectional
  - “Half-duplex”
  - Byte-stream
    - How does the reader know EOF?
  - IPC between related processes(`fork()`)
- Notes for correctly using pipe
  - Always close the un-used read/write end of a pipe



# Pipe (2/4)

## example5.c

```
#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>
#include<fcntl.h>

int main(){

    int pipe_fd[2];
    int bytes_read;
    int total_bytes_read = 0;
    char readbuf[1024];

    /* create a pipe */
    int ret = pipe(pipe_fd);
    if(ret < 0){
        printf("create pipe error\n");
        return 0;
    }

    /* write something into the pipe */
    write(pipe_fd[1], "-----\npipe test1\n-----\n", 25);
    write(pipe_fd[1], "pipe test2\n-----\n", 18);

    /* close the write end, EOF will be added into the pipe */
    close(pipe_fd[1]);

    /* read from the pipe into a buffer until there is no content in the pipe */
    while ((bytes_read = read(pipe_fd[0], &readbuf[total_bytes_read], 1024-total_bytes_read)) > 0){
        total_bytes_read += bytes_read;
    }

    /* print the content in the buffer */
    printf("total read %d\n",total_bytes_read);
    readbuf[total_bytes_read]='\0';
    printf("%s\n", readbuf);

    /* close read end of pipe */
    close(pipe_fd[0]);
}
```

## ➤ Basic usage

### ➤ Function prototype

➤ int pipe(int pipefd[2]);

### ➤ Pipe() creates a pipe, a unidirectional data channel that can be used for inter-process communication.

➤ pipefd[0]: read end of the pipe

➤ pipefd[1]: write end of the pipe

➤ On success, zero is returned. On error, -1 is returned



# Pipe (3/4)

```
#include <assert.h>
#include <signal.h>
#include <stdbool.h>
#include <stdio.h> /* perror */
#include <stdlib.h>
#include <sys/wait.h> /* wait, sleep */
#include <unistd.h> /* fork, write */

int main() {
    pid_t pid;
    int wait_return, status;
    int pipe_fd[2];
    int bytes_read;
    int total_bytes_read = 0;
    char readbuf[1024];

    /* create a pipe */
    int ret = pipe(pipe_fd);
    if(ret < 0){
        perror("create pipe error\n");
        return 0;
    }
}
```

## example6.c

```
pid = fork();
if(pid == -1){
    perror("fork");
    assert(false);
}
else if(pid == 0){
    /* close un-used write end of the pipe */
    close(pipe_fd[1]);
    /* read from parent process into the buffer*/
    while ((bytes_read = read(pipe_fd[0], &readbuf[total_bytes_read], 1024-1-
total_bytes_read)) > 0){
        total_bytes_read += bytes_read;
    }
    /* print the content in the buffer */
    readbuf[total_bytes_read]='\0';
    printf("[Message from parent] %s\n", readbuf);
    exit(0);
}
else{
    /* close un-used read end of the pipe */
    close(pipe_fd[0]);

    // sleep(5);

    /* write something into the pipe */
    write(pipe_fd[1], "Hello child process", 20);

    /* finished writing, close write end so child sees EOF */
    close(pipe_fd[1]);

    /* wait for the child process to terminate */
    while((wait_return = wait(&status)) > 0);
}

return 0;
}
```

## ➤ IPC

- Pipe is one of the things that will be inherited by child processes created with fork()
- Inter process communication(IPC) can be achieved with pipe
  - pipefd[0]: read end of the pipe
  - pipefd[1]: write end of the pipe
  - On success, zero is returned. On error, -1 is returned

# Pipe (4/4)

## example7.c

```
#include <assert.h>
#include <signal.h>
#include <stdio.h> /* perror */
#include <stdlib.h>
#include <sys/wait.h> /* wait, sleep */
#include <unistd.h> /* fork, write */

void signal_handler(int sig) {
    char s1[] = "SIGPIPE captured\n";
    char s2[] = "Unknown Signal\n";
    if (sig == SIGPIPE){
        write(STDOUT_FILENO, s1, sizeof(s1));
    }
    else{
        write(STDOUT_FILENO, s2, sizeof(s2));
    }
}

int main() {
    pid_t pid = getpid();
    int pipe_fd[2];

    /* create a pipe */
    int ret = pipe(pipe_fd);
    if(ret < 0){
        printf("create pipe error\n");
        return 0;
    }

    /* create a signal handler to capture signal pipe */
    signal(SIGPIPE, signal_handler);

    /* close all read end of the pipe */
    close(pipe_fd[0]);

    /* writer something to the pipe */
    write(pipe_fd[1], "test content", 12);

    return 0;
}
```

## ➤ Pipe w/ read/write

### ➤ example6.c

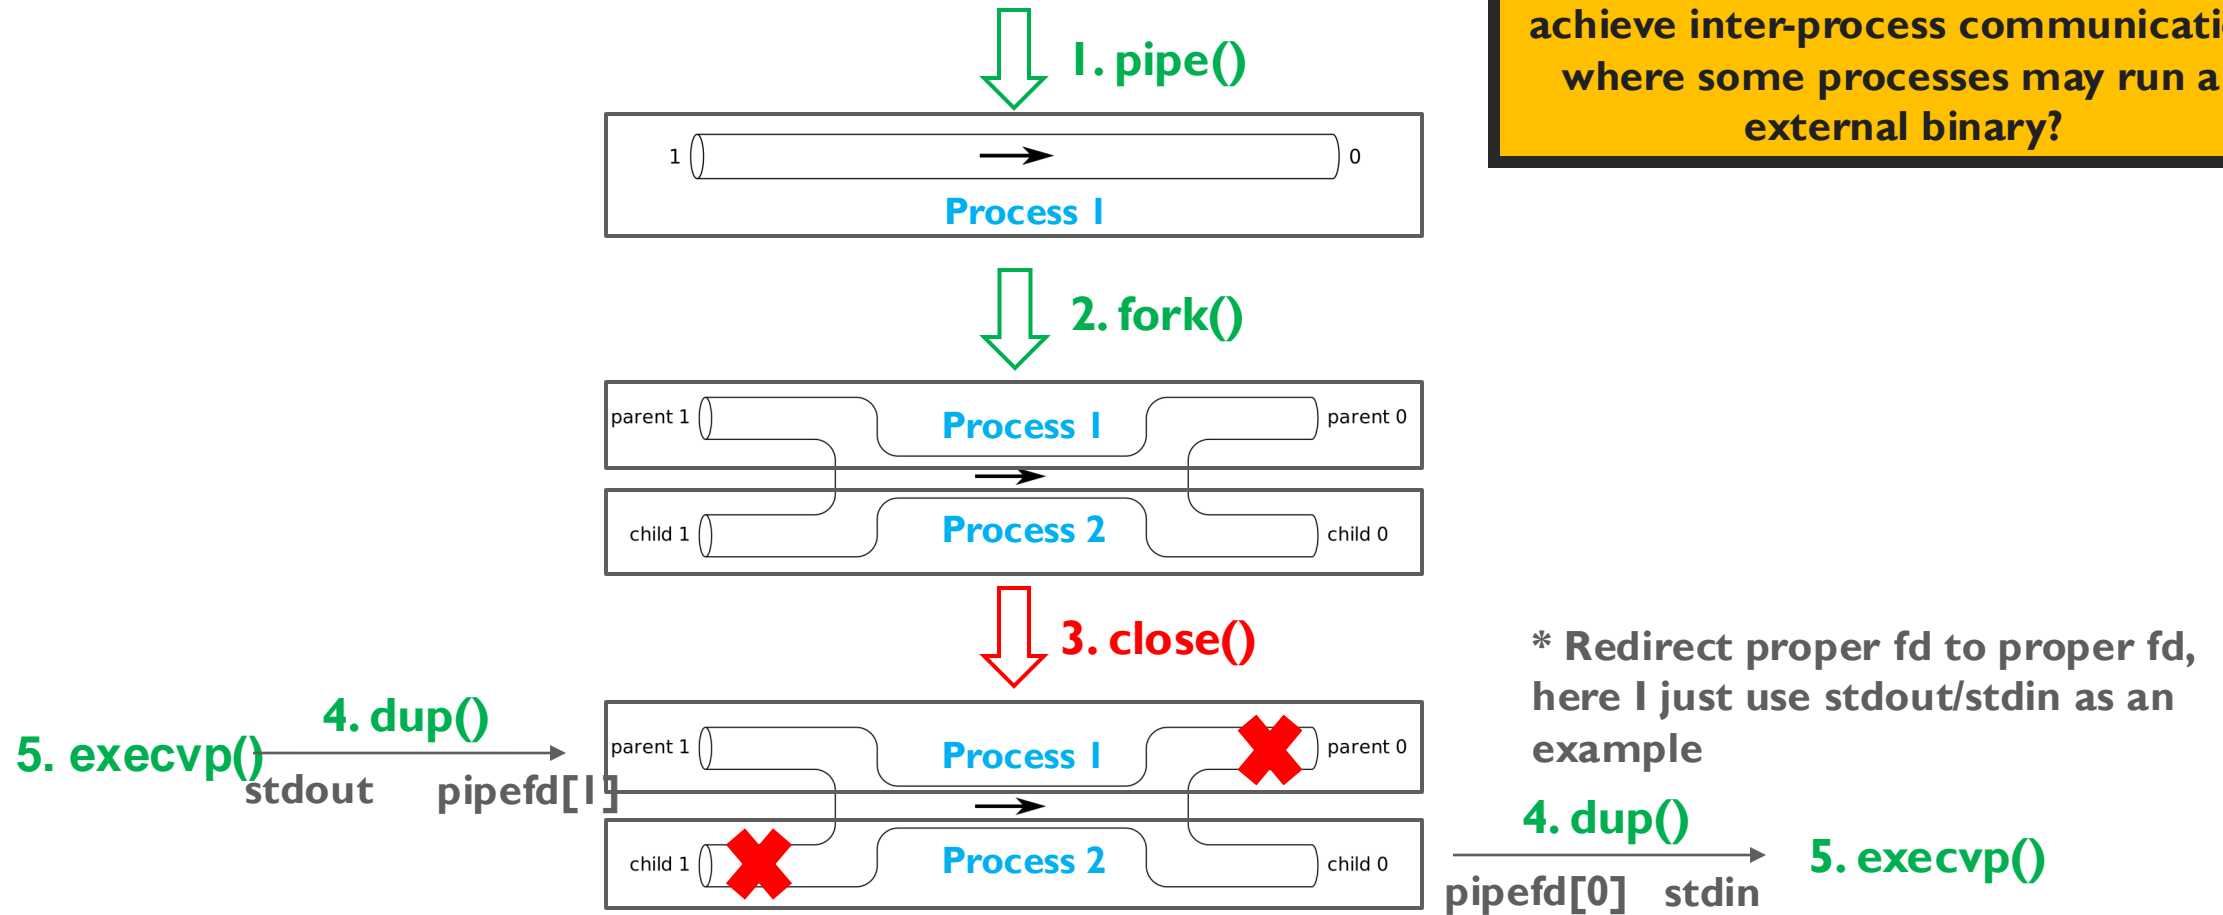
- If all file descriptors referring to the write end of a pipe have been closed
  - Read from pipe will see end-of-file(EOF)
- If a process attempts to read from an empty pipe
  - Read will block until data is available

### ➤ example7.c

- If all file descriptors referring to the read end of a pipe have been closed
  - Write to the pipe will generate SIGPIPE for the process
- If a process attempts to write to a full pipe
  - Write blocks until sufficient data has been read from the pipe to allow the write to complete

# Summary (I/I)

How to use the above systems calls to achieve inter-process communication, where some processes may run an external binary?



**Q & A**